

COL380

Introduction to  
Parallel & Distributed Programming  
2-0-2

Subodh Kumar

# Agenda

- Course structure and policies
- Review of basic concepts
  - ➔ OS, Compilers, Architecture

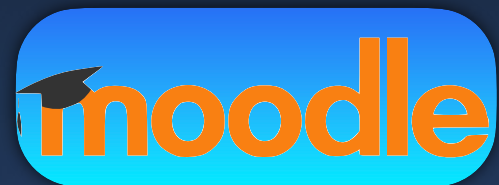
- COL331 is a ~~pre-~~ co-requisite
- Those without COL331 will be de-registered at the end of the add-drop period

# Resources

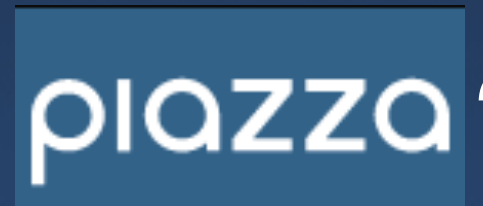


- <http://www.cse.iitd.ac.in/~subodh/courses/COL380>

→ Persistent Info: Policies, Resources, Links, Slides, Textbook



- Assignments, Test, Discussion



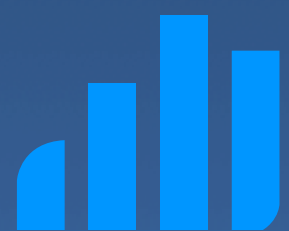
- Course discussion



- Urgent announcements

[subodh@cse.iitd.ac.in](mailto:subodh@cse.iitd.ac.in)

Subject: [COL380] ...



- Grading

- Lectures
  - Tuesday and Friday
- Mondays for doubts, Q&A
  - And announced quizzes



- Assignments: 29.

- \* Do not share code
  - \* Do not discuss, except with me or TA

- A0: 4, A1-A3: 5 each, A4: 10 Two deadlines, resp. correct and performant

- Marks for performance (only if correct); >100% possible

- Checked for similarity; some will give **viva**

Similarity 1 ⇒

0 mark + 1 letter-grade penalty

- Programming Test: 6.

Similarity 2 ⇒

F grade + Disciplinary com.

- Quiz(zes): 10.

- Announced / Unannounced in class

- Minor: 25. Major 30.

- Objective + subjective

- Late Assignment Submission:
  - 0.25 marks/6H of delay or part thereof; up to 3 days of delay
- 100% attendance is required
  - <75% in any month  $\Rightarrow$  take a **viva** on the material covered that month
    - ▶ Range of marks in the viva: -2.5 to 0
  - 75% pre-minor required for reminor, 75% in semester for remajor
- Audit Policy:
  - No audits

## Other Policies

- **Late Assignment Submission:**

- 0.25 marks/6H of delay or part thereof; up to 3 days of delay

### VIVA

- **100% attendance** Usually on Monday 12-1 (Keep slot free.)

- <75% in any month → take a viva on the material covered that month  
Viva call by email before that Monday at 8am  
If viva not on Monday, email will be ~24 hours in advance

- ▶ Range of marks in the viva: -2.5 to 0

If you fail to show up, you will get the minimum marks

- 75% pre-minor required for re-minor, 75% in semester for repeat
- If you are sick, send a request to reschedule BEFORE the viva

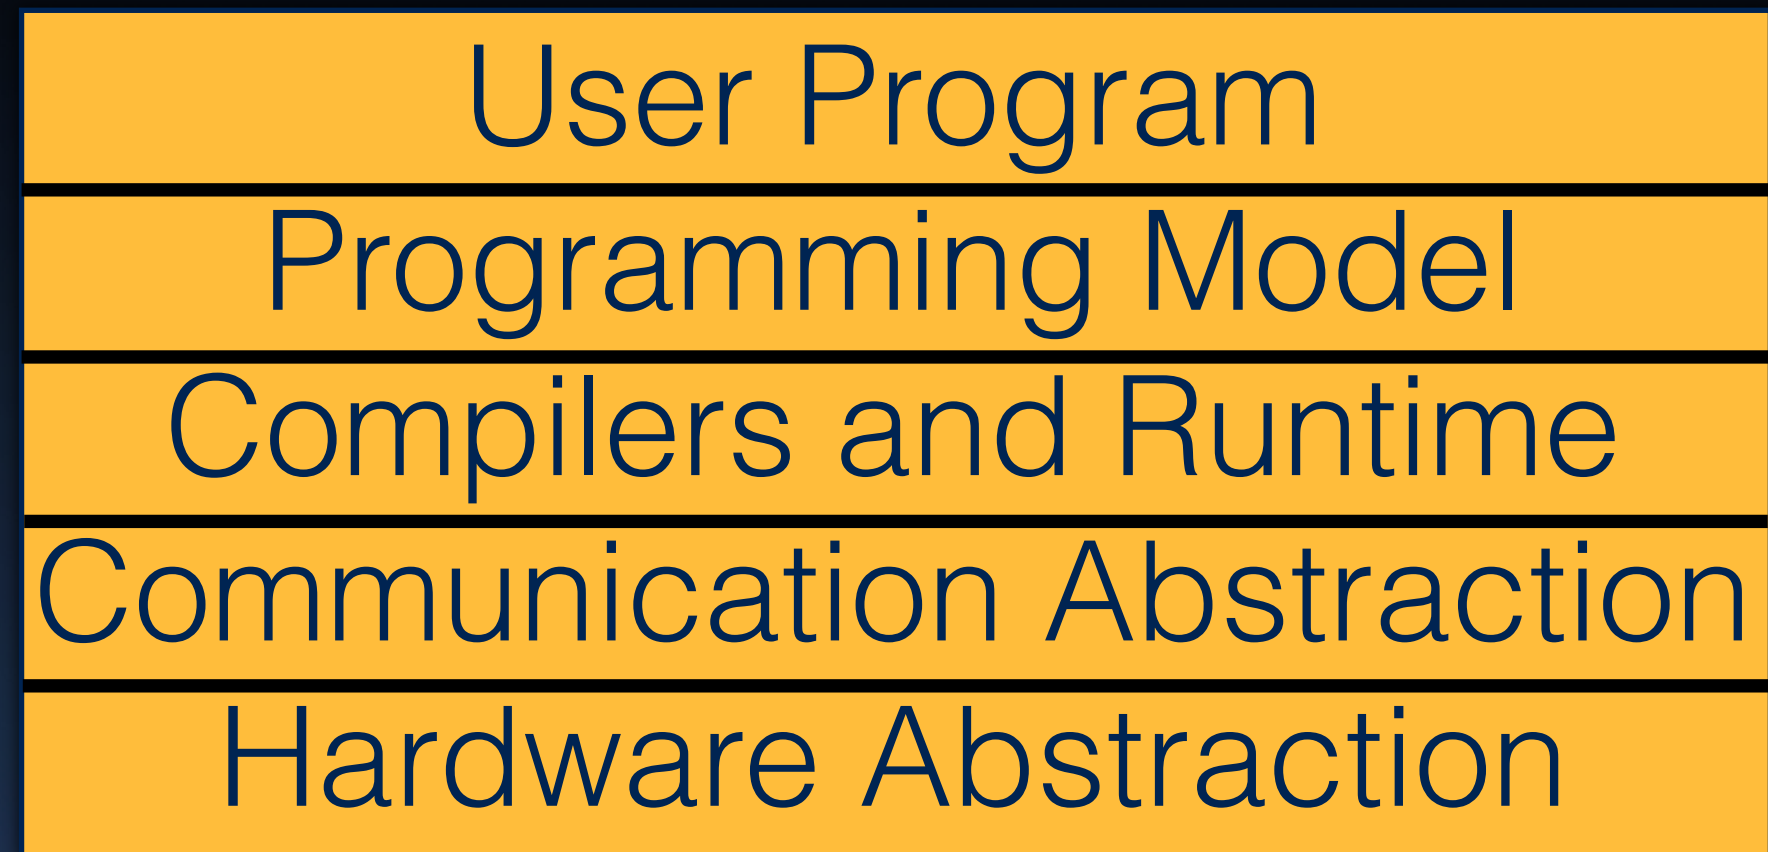
- **Audit Policy:**

(In case of such request, it is YOUR responsibility to find an alternate time within 1 week of the original time.)

- No audits



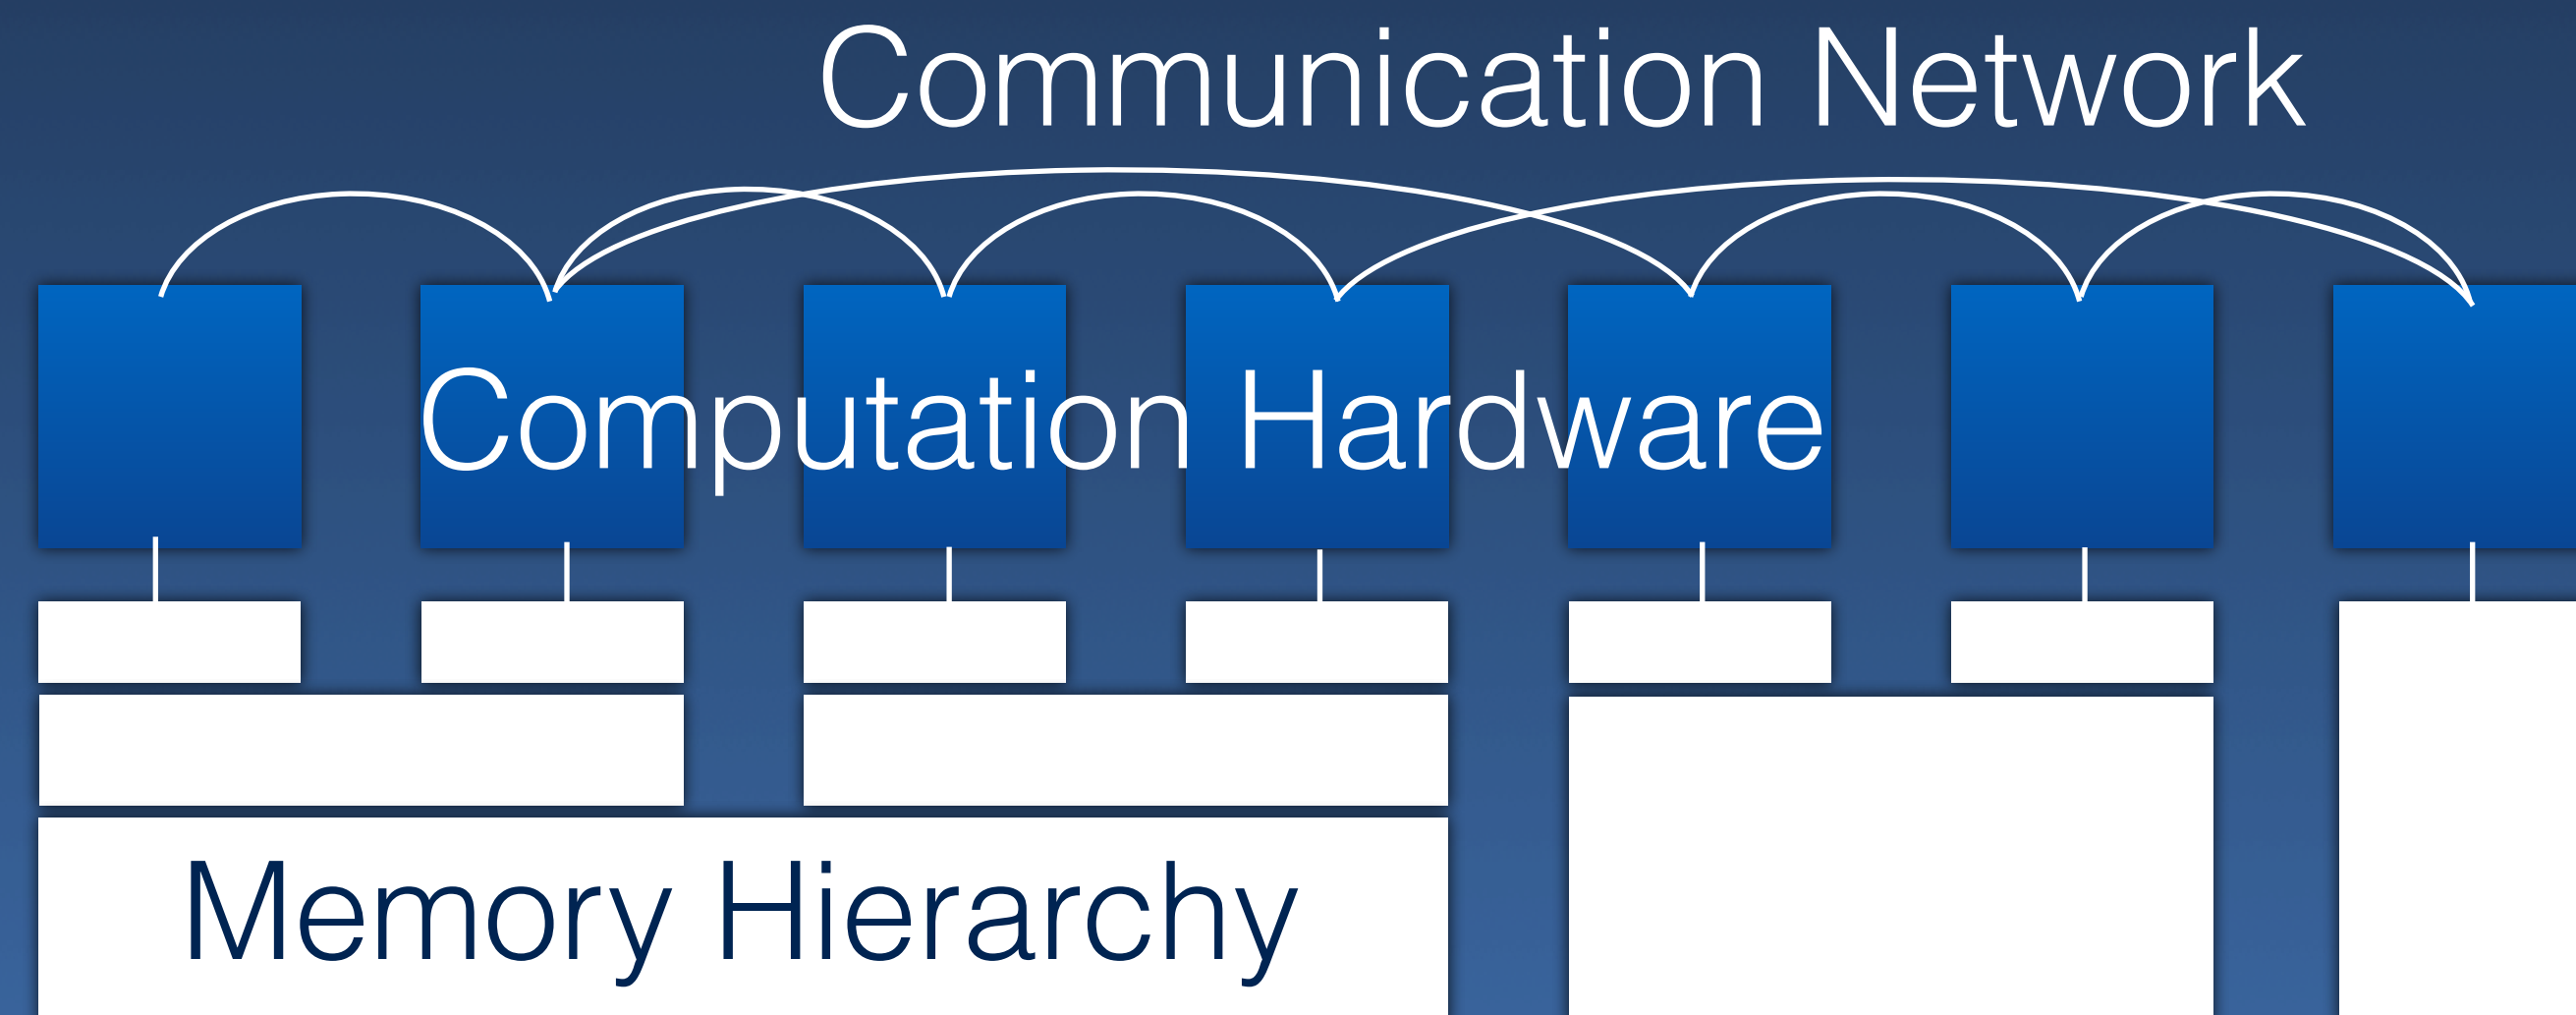
## About



- Decide what to do
- Where
- When
- How to ask the system to do it
- Estimate and Measure Efficiency

- Algorithms
- Programs
- System S/W
- Hardware

Synchronization  
Communication  
Load distribution  
Latency  
Bandwidth  
Scaling



# Course Content

Introduction to concurrency, Race conditions, Atomicity, Semantics of concurrent programs, Examples of distributed algorithms, Client-Server paradigm

Parallel architecture, Flynn's classification

Shared-memory programming with reference to memory consistency, cache in-coherence, false sharing and mutual exclusion

Message passing, High level and collective constructs, Point-to-point communication, multicast and broadcast, Blocking versus non-blocking styles for communication, Message buffering

Theoretical models of parallel computation and algorithm analysis, Examples of reduction, prefix-sum

Performance metrics: Time, work, Scalability.

Task/Communication Dependence graphs, Task decomposition, Data-parallel decomposition, Pipelining

Synchronization, barriers, Progress, Livelock/Deadlock

## Learning Goals

- Write scalable and efficient parallel programs
  - ➔ OpenMP, MPI, Cuda, Development tools
  - ➔ Understand the issues in tool design and implementation
  - ➔ Profile and Debug
- Understand, measure, predict and analyze parallel performance
- Examples of parallel and distributed algorithms and data structures
- Understand parallelism in I/O and memory
- Understand nomenclature, literature, documentation



## Learning Goals

- Write scalable and efficient parallel programs
  - ➔ OpenMP, MPI, Cuda, Development tools
  - ➔ Understand the issues in tool design and implementation
  - ➔ Profile and Debug
- Understand, measure, predict and analyze parallel performance
- Examples of parallel and distributed algorithms and data structures
- Understand parallelism in I/O and memory
- Understand nomenclature, literature, documentation

### Need

- strong C++ skills
- OS and Architecture concepts
- Background in Algorithms



## Keys to success

- Be regular (Slides only have topics)



- Read



- Program all assignments yourself



→ Be curious. Try out 'what-if' variations and see what happens

- Talk to the instructor



## First Parallel Program

### Instruction level Parallelism

```
float dot(float *a, float *b, int n)
{
    int i;
    float s0=0, s1=0, s2=0, s3=0;
    for(i=0; i<n/4*4; i+=4)
        s0 += a[i]*b[i];
        s1 += a[i+1]*b[i+1];
        s2 += a[i+2]*b[i+2];
        s3 += a[i+3]*b[i+3];
    for(; i<n; i++)
        s0 += a[i]*b[i];
    return s0+s1+s2+s3;
}
```

unrolled loop

Independent statements

Architecture can issue multiple instructions

Architecture can prefetch data

Compilers also re-order instructions

- Process & thread
  - ➔ Scheduling, Context-switching and concurrency
  - ➔ User space, Kernel space
  - ➔ System calls
- Address space & Name space
- Virtual Memory
  - ➔ Caches
- Synchronization

- ◆ Shell (bash)
- ◆ time
- ◆ PMU (perf)
- ◆ Context switch
  - ◆ and scheduling
- ◆ System calls
- ◆ Interrupts