

COL380

Introduction to
Parallel & Distributed Programming

OpenMP Threads & Tasks


- Tasks are a software construct
 - ➔ Have execution context: stack, heap, PC..
 - ➔ To be executed by some “device”
- In OpenMP tasks are not directly scheduled on hardware devices
 - ➔ It instead uses the intermediate construct “Thread”
 - ➔ Threads are (can be) scheduled by the OS
- User-space code schedules/assigns tasks to threads

```
#pragma omp parallel  
#pragma omp single  
#pragma omp task  
    processAB(a, b, n);
```

Sharing and Racing

```
void loop(int &a, int n)
{
    for(int i=0; i<n; i += 3)
        a++;
}

..
int a = 0, n = 10;
#pragma omp parallel
    loop(a, n);
```

- Operation are **not instantaneous**: 
- Such an operation can becomes visible to different threads at different times
- apparent order of operations may not be *consistent*
- Even *consistent* operations can be *concurrent* (vary from execution to execution)
- Program must remain correct no matter the order

Data race: Concurrent RW or WW operations

Race condition: If variable order of concurrent operations affects correctness

```
void loop(int &a, int n)
{
    for(int i=0; i<n; i += 3)
        a++;
}
(Rd a; Add1; Wr a)

..
int a = 0, n = 10;
#pragma omp parallel
    loop(a, n);
```

```
    movl    (%rdi), %eax
    addl    $1, %eax
    movl    $0, %edx
.L3: movl    %eax, %ecx
    addl    $3, %edx
    addl    $1, %eax
    cmpl    %edx, %esi
    jg      .L3
    movl    %ecx, (%rdi)
```



```
void loop(volatile int &a, int n)
{
    for(int i=0; i<n; i += 3)
        a++;
}
```

(Rd a; Add1; Wr a)

```
..
int a = 0, n = 10;
#pragma omp parallel
    loop(a, n);
```

```
    movl    (%rdi), %eax
    addl    $1, %eax
    movl    $0, %edx
.L3: movl    %eax, %ecx
    addl    $3, %edx
    addl    $1, %eax
    cmpl    %edx, %esi
    jg      .L3
    movl    %ecx, (%rdi)
```

```
.L3: movl    (%rdi), %eax
    addl    $1, %eax
    movl    %eax, (%rdi)
    addl    $3, %edx
    cmpl    %edx, %esi
    jg      .L3
    ret
```

Share Occasionally

```
void loop(volatile int &a, int n)
{
    for(int i=0; i<n; i += 3)
        a++;
}
```

(Rd a; Add1; Wr a)

```
..
int a = 0, n = 10;
#pragma omp parallel
    loop(a, n);
```

```
    movl    (%rdi), %eax
    addl    $1, %eax
    movl    $0, %edx
.L3: movl    %eax, %ecx
    addl    $3, %edx
    addl    $1, %eax
    cmpl    %edx, %esi
    jg      .L3
    movl    %ecx, (%rdi)
```

```
.L3: movl    (%rdi), %eax
    addl    $1, %eax
    movl    %eax, (%rdi)
    addl    $3, %edx
    cmpl    %edx, %esi
    jg      .L3
    ret
```

var = val need not be “seen”
or, need not be “atomically” seen

Atomic = “Not divisible”
(with respect to a set of operations)

Memory Model

@atom

Address space

Memory Model

- Register size?
- Limit on the number of readers and writers?
- Asynchronous reads and writes
 - ➔ Global knowledge of 'time?'



Memory Model

- Register size?
- Limit on the number of readers and writers?
- Asynchronous reads and writes
 - ➔ Global knowledge of 'time?'
- * What is the value read? (Correctness)

“Most recent Write”



Memory Model

- Register size?
- Limit on the number of readers and writers?
- Asynchronous reads and writes
 - Global knowledge of 'time?'
- * What is the value read? (Correctness)
- Increasing power \Rightarrow reducing performance



Memory Model

- Register size?
- Limit on the number of readers and writers?
- Asynchronous reads and writes
 - ➔ Global knowledge of 'time?'
- * What is the value read? (Correctness)
- Increasing power \Rightarrow reducing performance



Same ideas apply to higher level data structures

thread A: —Pu-3—Po-2→

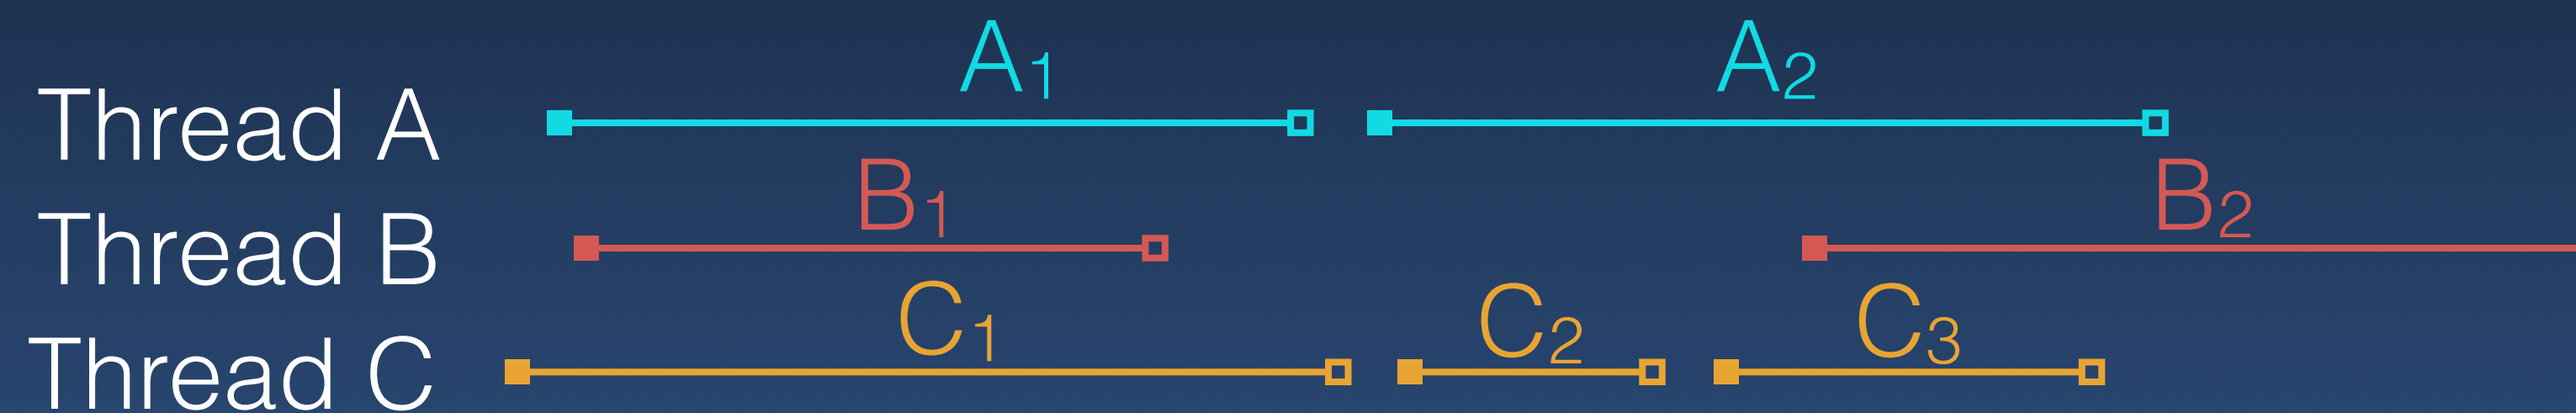
thread B: —Pu-2—Po-3→

Stack

- Operations appear to take global effect at some instant

→ between their start and end

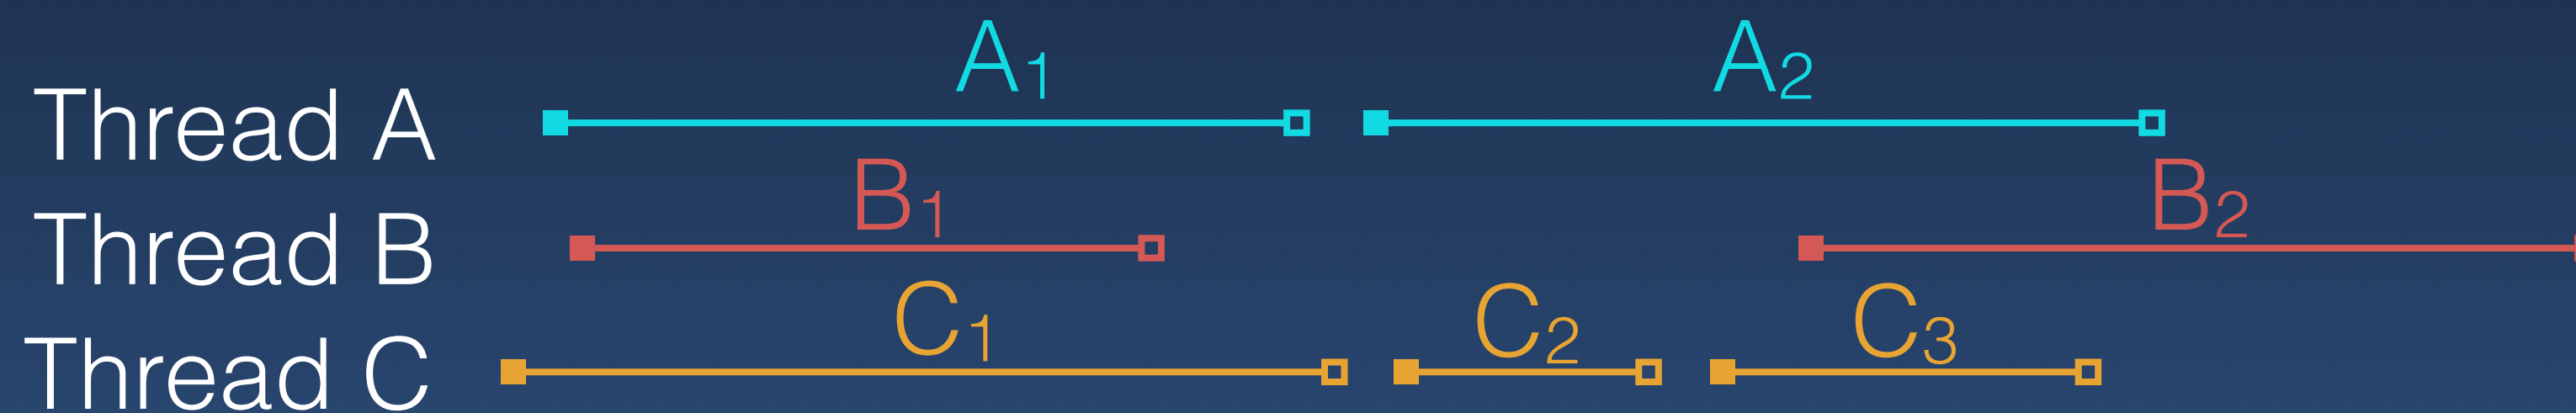
→ If there is no overlap, ordering is well defined
(knowledge of global time)



- Operations appear to take global effect at some instant

→ between their start and end

→ If there is no overlap, ordering is well defined
(knowledge of global time)

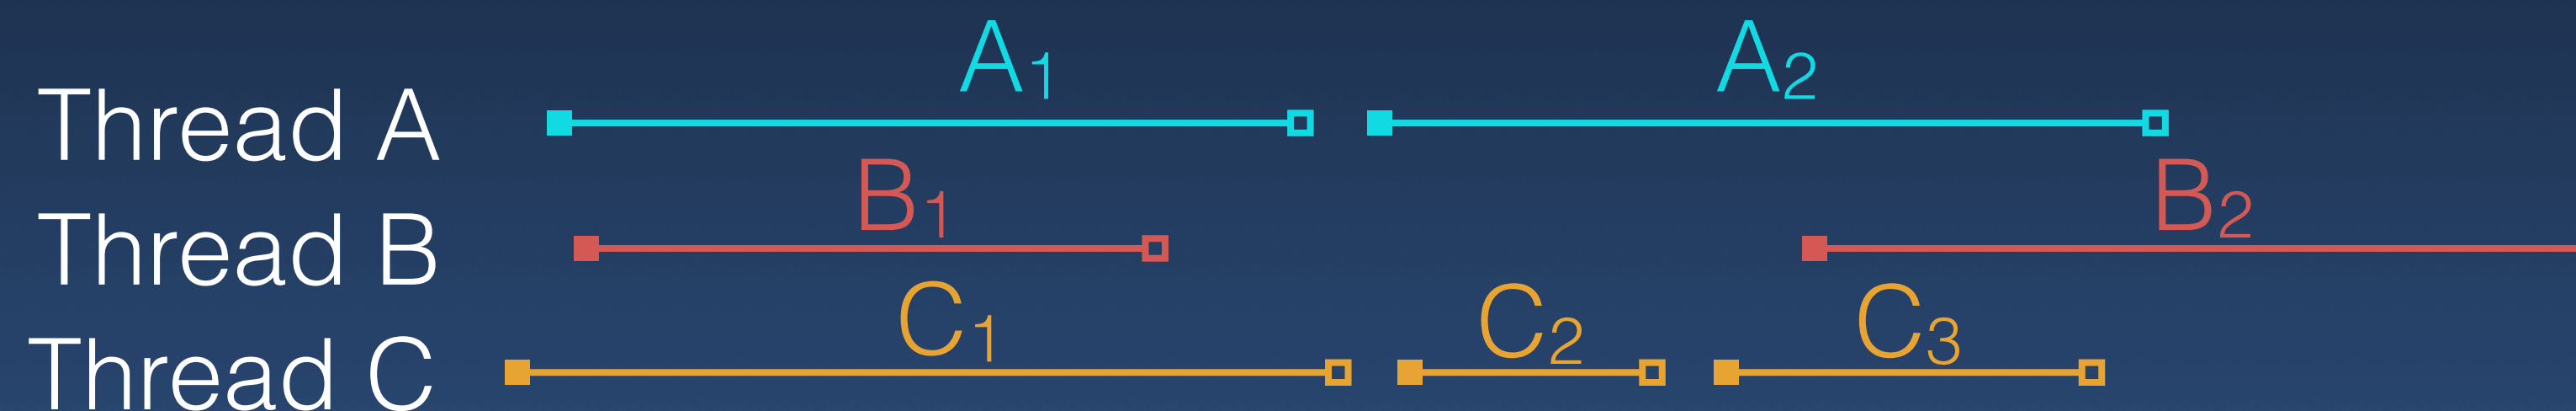


Equivalent sequential history S exists: A₁ C₁ B₁ C₂ B₂ C₃ A₂

- Operations appear to take global effect at some instant

→ between their start and end

→ If there is no overlap, ordering is well defined
(knowledge of global time)



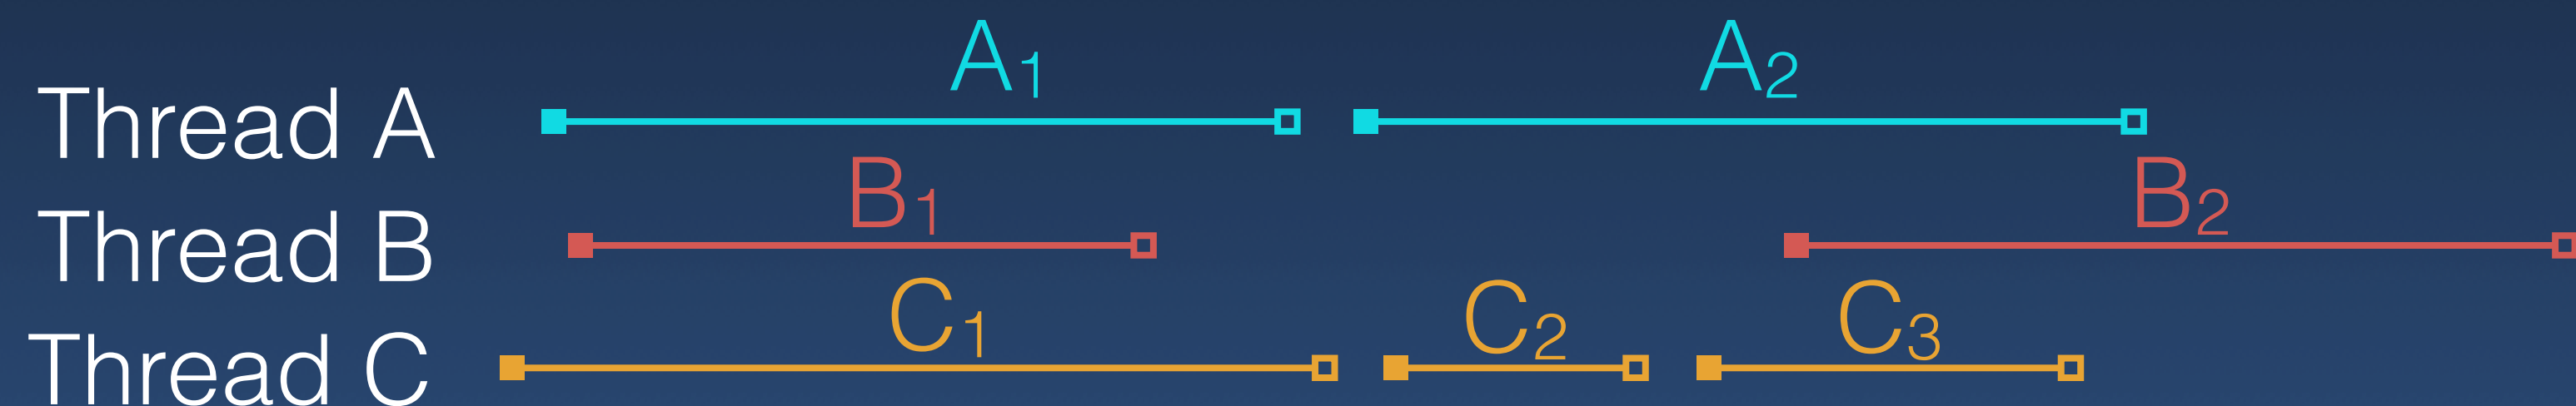
A₁, C₁ overlap

Equivalent sequential history S exists: C₁ A₁ B₁ C₂ B₂ C₃ A₂

- Operations appear to take global effect at some instant

→ between their start and end

→ If there is no overlap, ordering is well defined
(knowledge of global time)



A₁, C₁ overlap



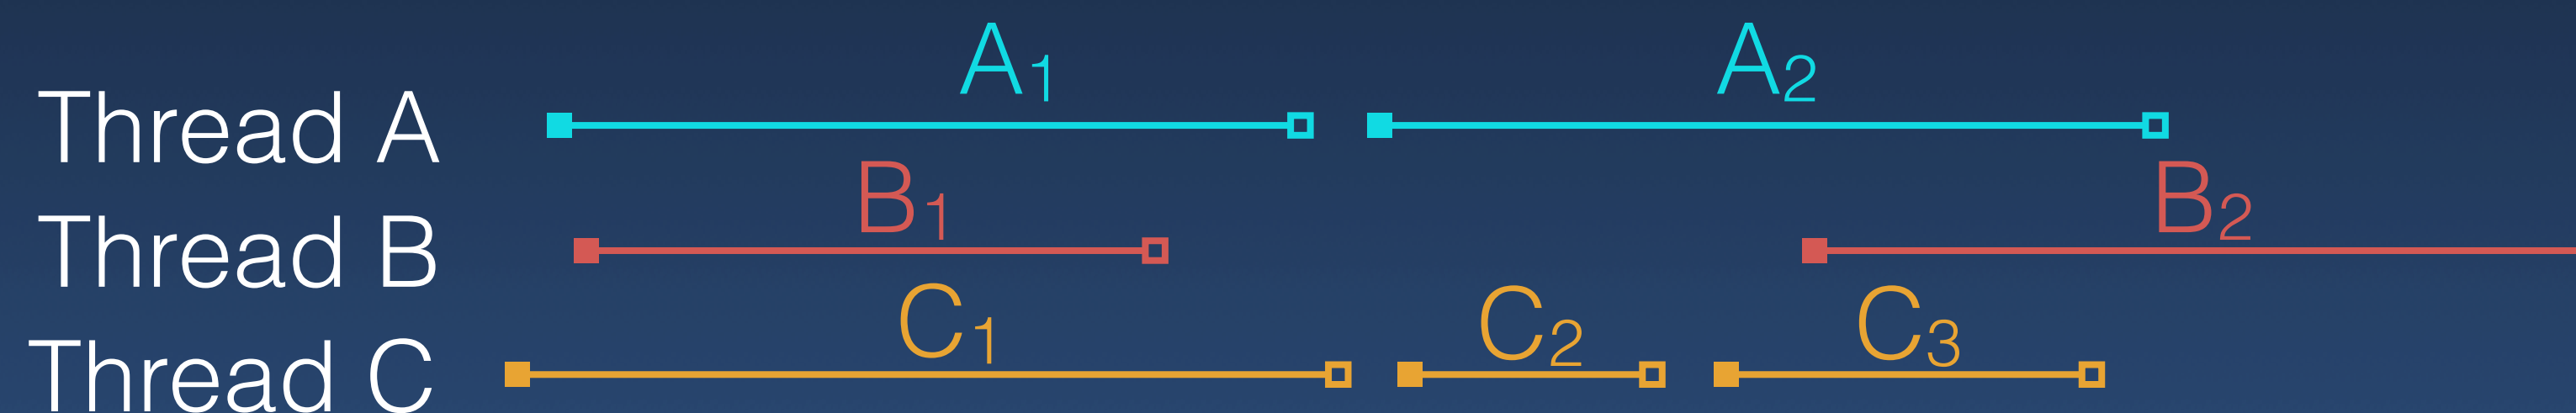
B₁, C₂ do not overlap

Equivalent sequential history S exists: C₁ A₁ B₁ C₂ B₂ C₃ A₂

- Operations appear to take global effect at some instant

→ between their start and end

→ If there is no overlap, ordering is well defined
(knowledge of global time)



Equivalent sequential history S exists: C₁ A₁ B₁ C₂ B₂ C₃ A₂

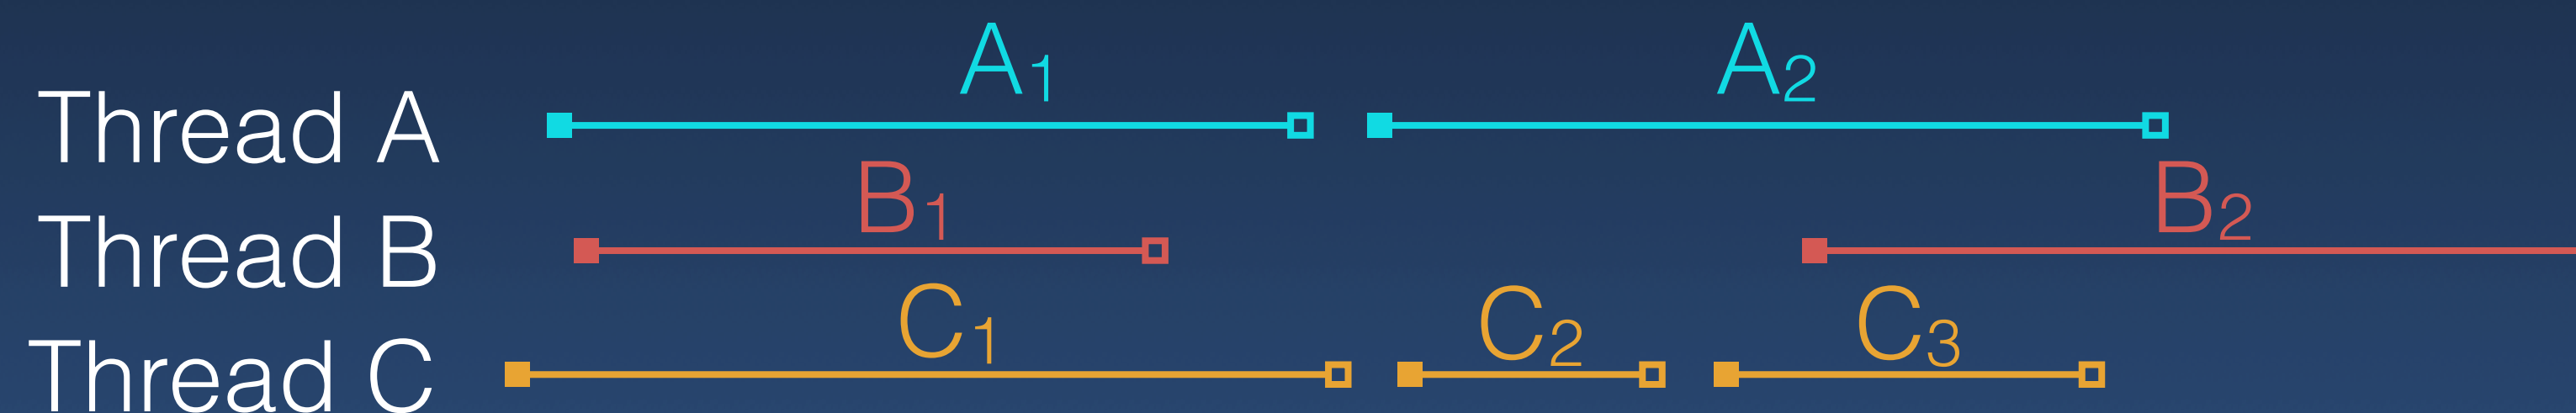
1. Got the result S would get
2. No thread's history is violated in S
3. Non-overlapping operations retain order in S

Linearizable

- Operations appear to take global effect at some instant

→ between their start and end

→ If there is no overlap, ordering is well defined
(knowledge of global time)



Equivalent sequential history S exists:

1. Got the result S would get
2. No thread's history is violated in S
3. Non-overlapping operations retain order in S

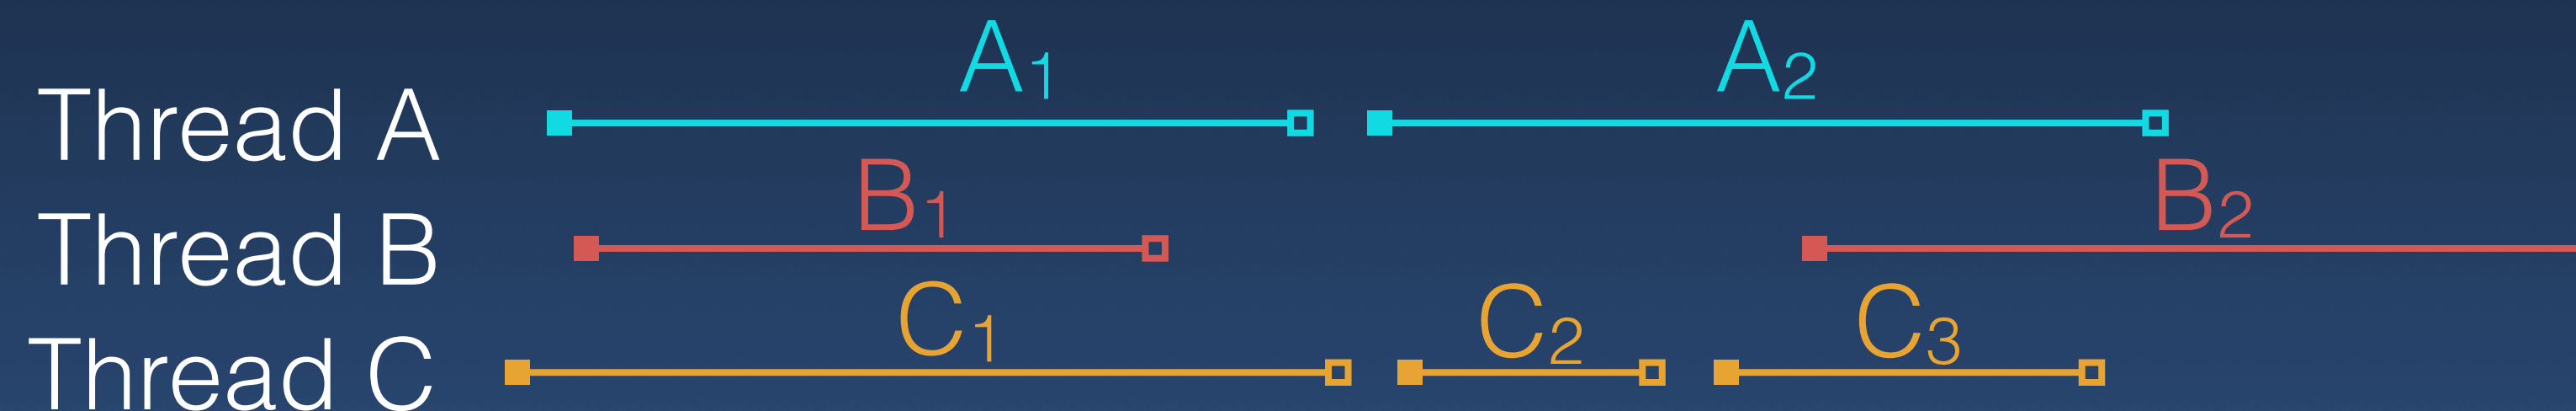
Linearizability is Composable

Linearizable

- Operations appear to take global effect at some instant

→ between their start and end

→ If there is no overlap, ordering is well defined
(knowledge of global time)



Equivalent sequential history S exists: 

Linearizability is Composible

1. Got the result S would get
2. No thread's history is violated in S
3. Non-overlapping operations retain order in S

What if we cannot tell what overlaps what

Linearizable

- Operations appear to take global effect at some instant

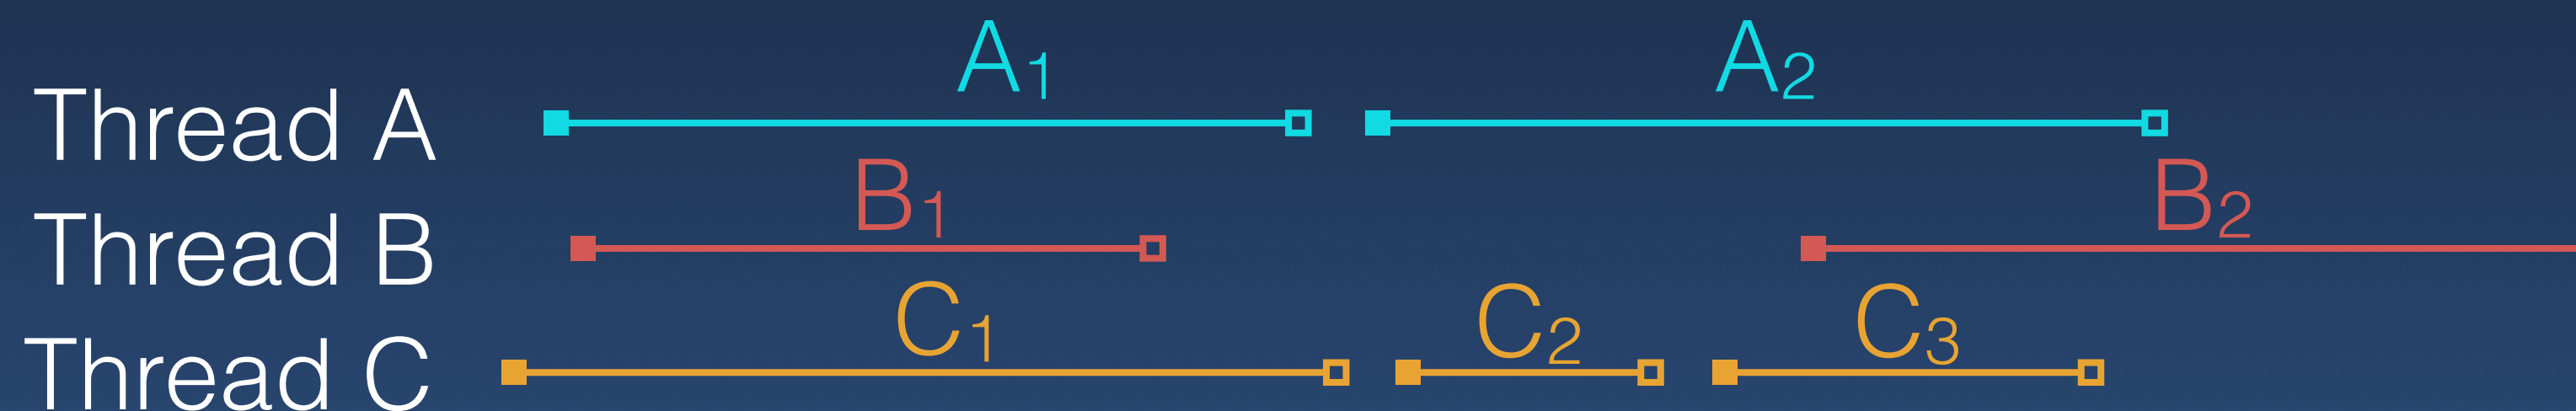
→ between their start and end

→ If there is no overlap ordering is well defined
(knowledge of global time)

Strict Consistency:

Operations are instantaneous

👉 Directly gives sequential order



Equivalent sequential history S exists:
C₁ A₁ B₁ C₂ B₂ C₃ A₂

1. Got the result S would get
2. No thread's history is violated in S
3. Non-overlapping operations retain order in S

Linearizability is Composable

What if we cannot tell what overlaps what

Types of Registers

- Linearizable registers

- ➔ SRSW 1-bit safe register

Can build shared registers with high guarantee using ones with low guarantee

- ➔ MRMW n-bit atomic linearizable register

Can enforce order between **read/write** using shared registers (general synchronization)

- Sequentially consistent registers

- Causally consistent registers

- FIFO consistent registers

- Weakly consistent registers