


COL380

Introduction to  
Parallel & Distributed Programming

“A multiprocessor is **sequentially consistent** if the result of any execution is the same as if the operations of all the processors were executed in **some sequential order**, and the operations of each individual processor appear in this sequence **in the order specified by its program**.” [Lamport, 1979]


Weaker than Linearizability  
Does not depend on global time  
(Still not efficient to guarantee.)

# Sequentially Consistent



- No global notion of time
  - ➔ Only consistent Order


# Sequentially Consistent



- No global notion of time
  - ➡ Only consistent Order



# Sequentially Consistent



- No global notion of time

→ Only consistent Order

Thread A:  $A_1 \longrightarrow A_2$

Thread B:  $B_1 \longrightarrow B_2$

Thread C:  $C_1 \longrightarrow C_2 \longrightarrow C_3$

## Sequential History

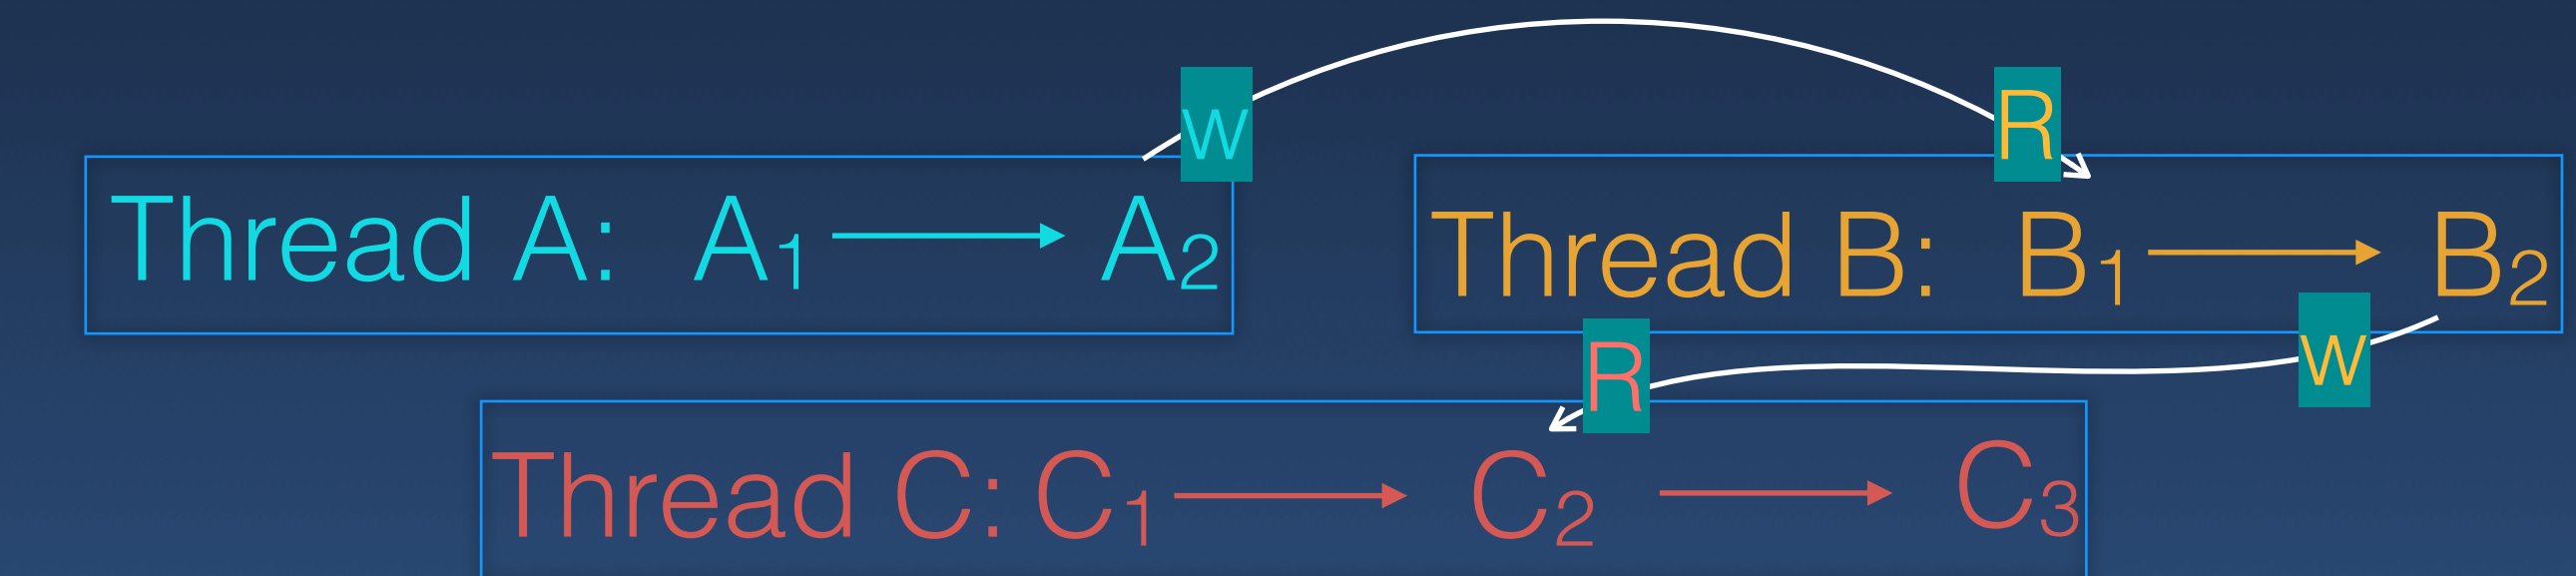
$A_1 \rightarrow A_2$   $C_1 \rightarrow C_2 \rightarrow C_3$   $B_1 \rightarrow B_2$

$A_1$   $C_1$   $A_2$   $B_1 \rightarrow B_2$   $C_2 \rightarrow C_3$


# Sequentially Consistent




- No global notion of time  
→ Only consistent Order




## Sequential History




# Sequentially Consistent



- No global notion of time  
→ Only consistent Order



## Sequential History

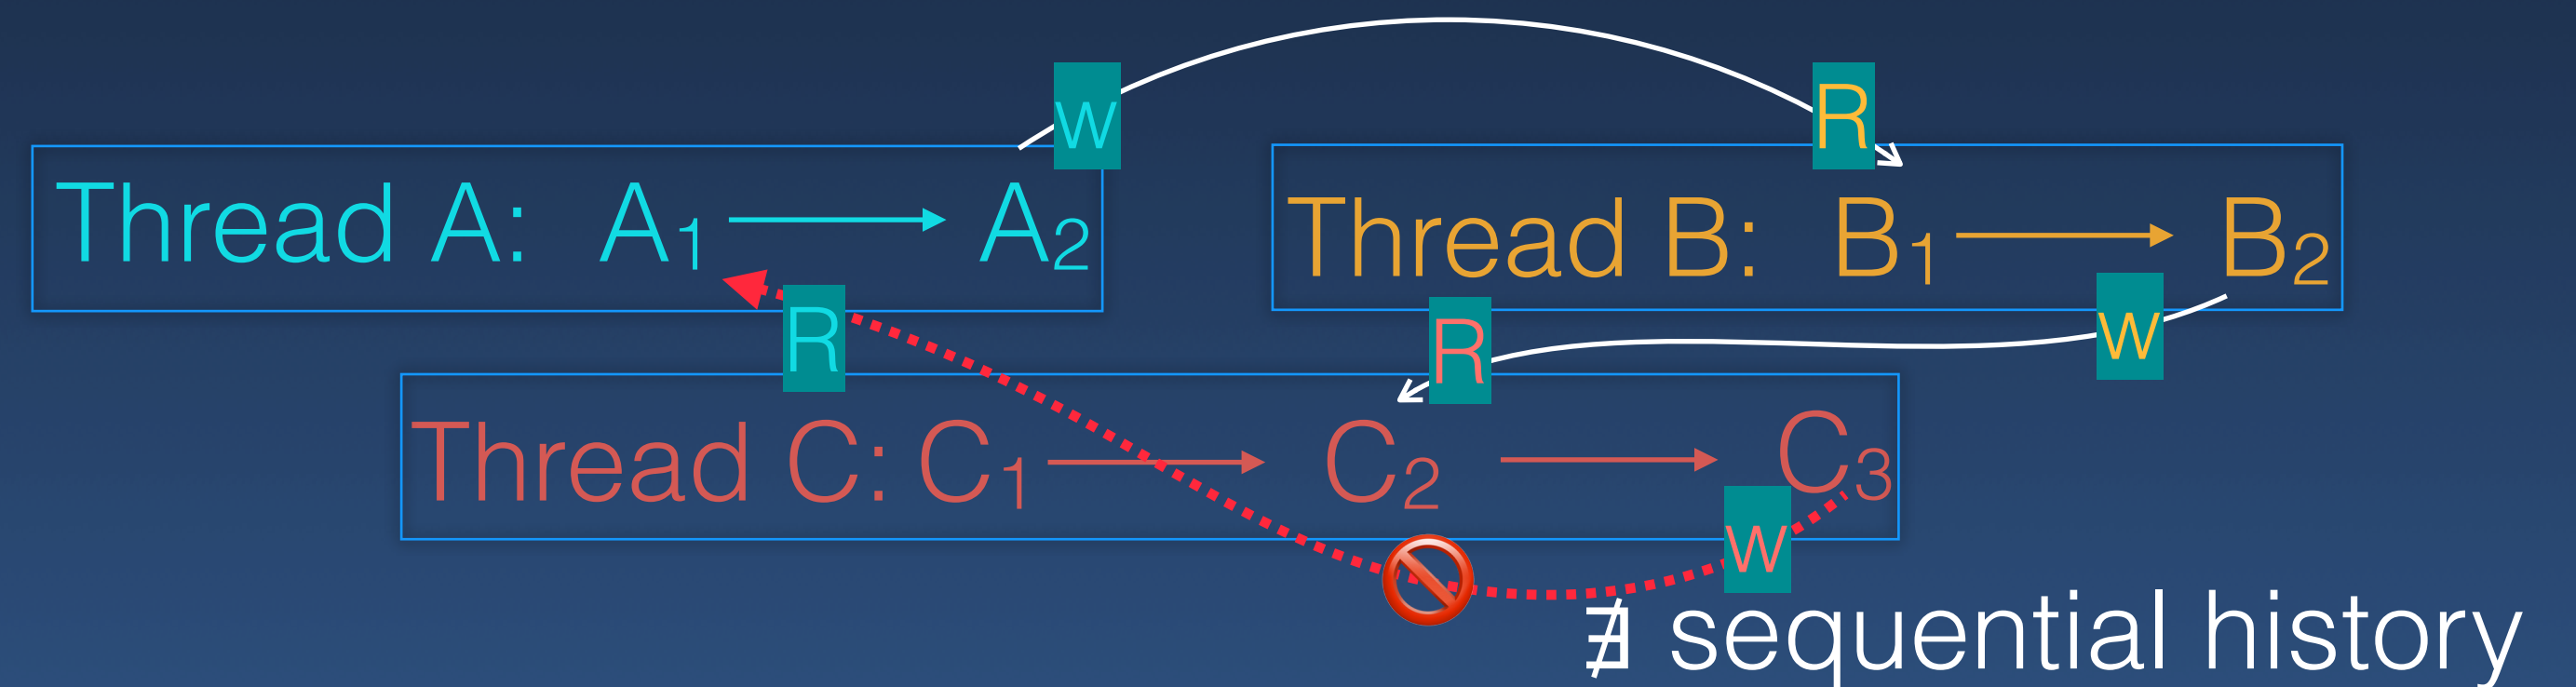





# Sequentially Consistent



- No global notion of time  
→ Only consistent Order



Sequential History

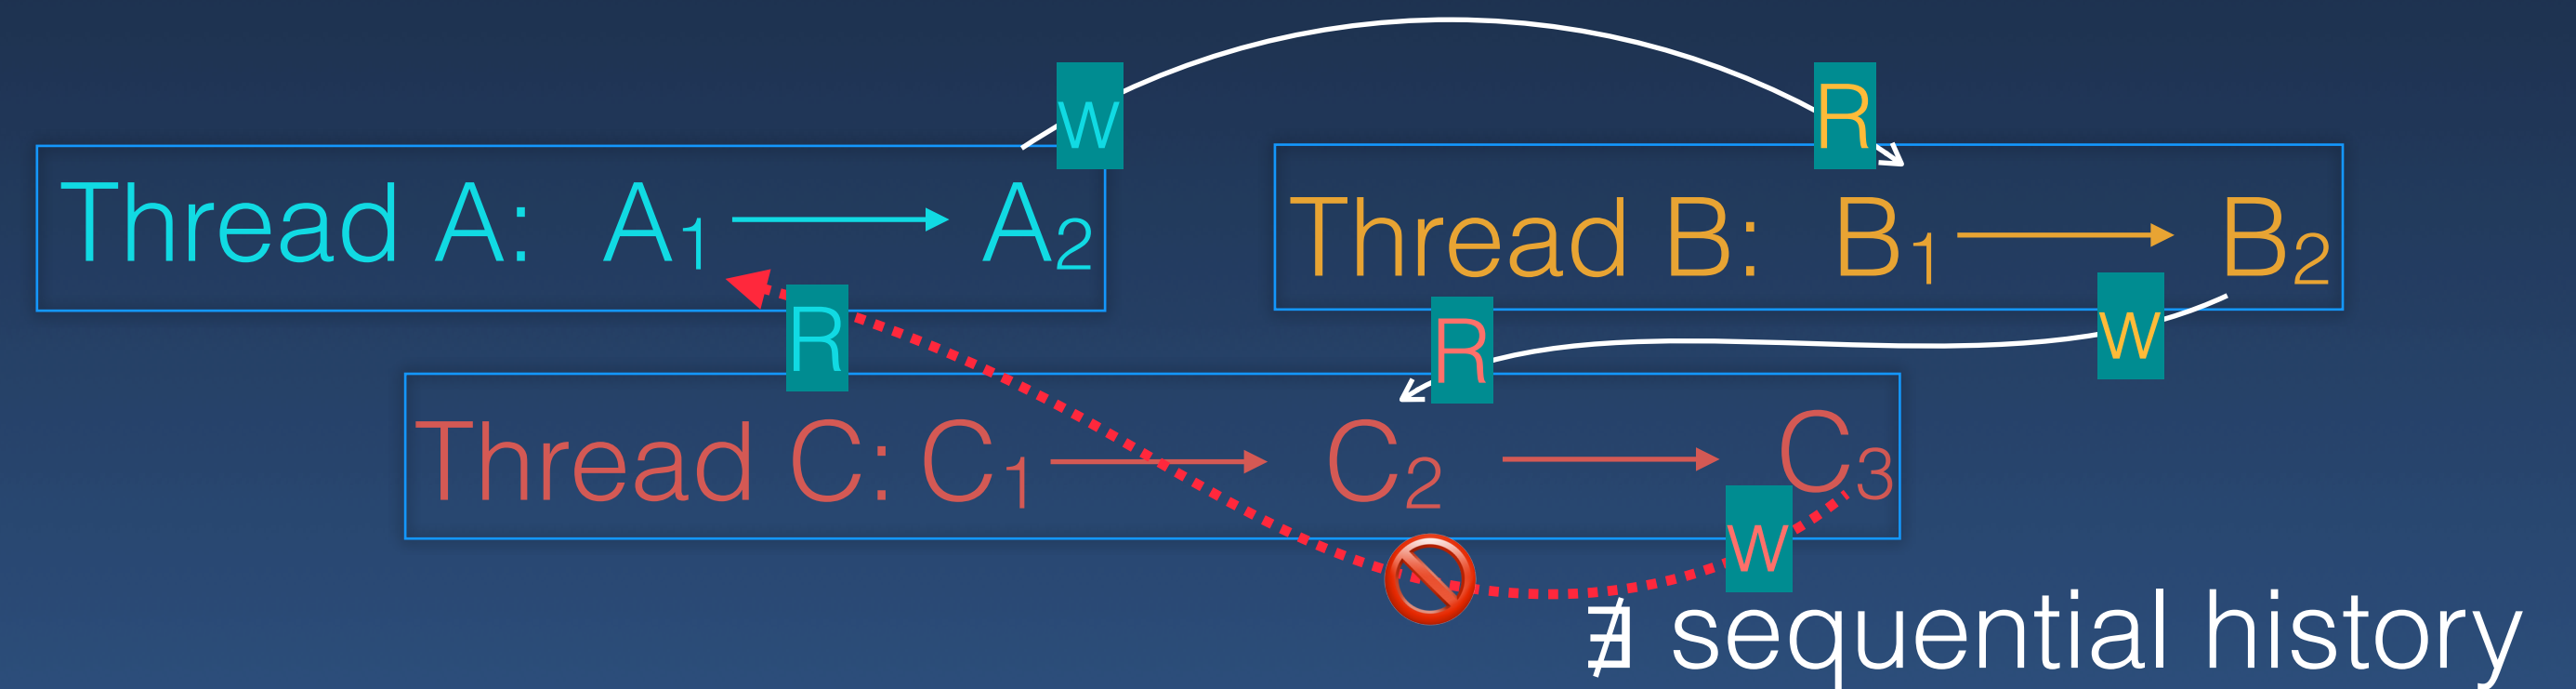




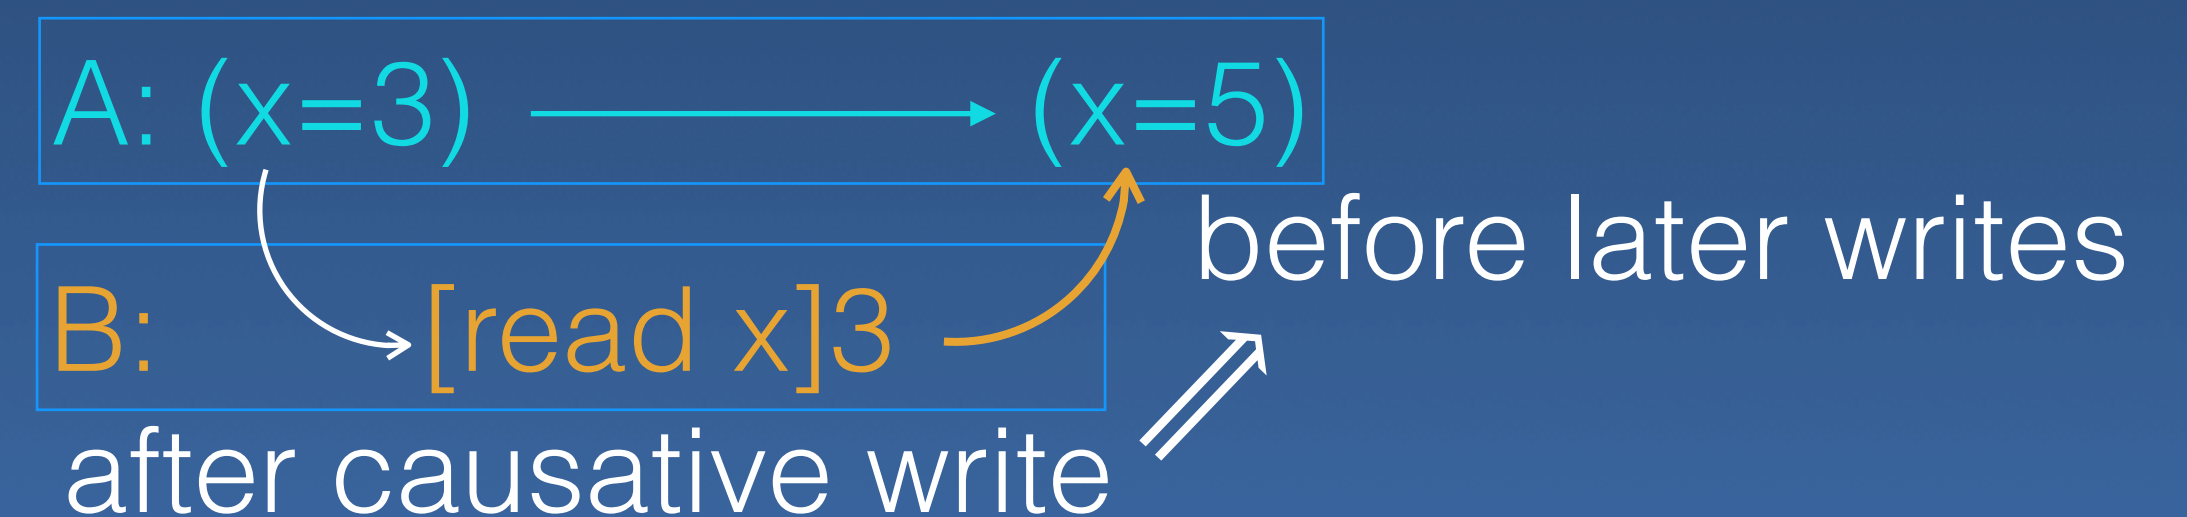

# Sequentially Consistent



- No global notion of time
  - Only consistent Order



Sequential History



# Applying Sequential Consistency

- Threads always see values written by some thread
  - ➔ No garbage (update is atomic)
- The value seen is constrained by thread-order
  - ➔ for every thread

initially: ready=0, data=0

thread P

thread C

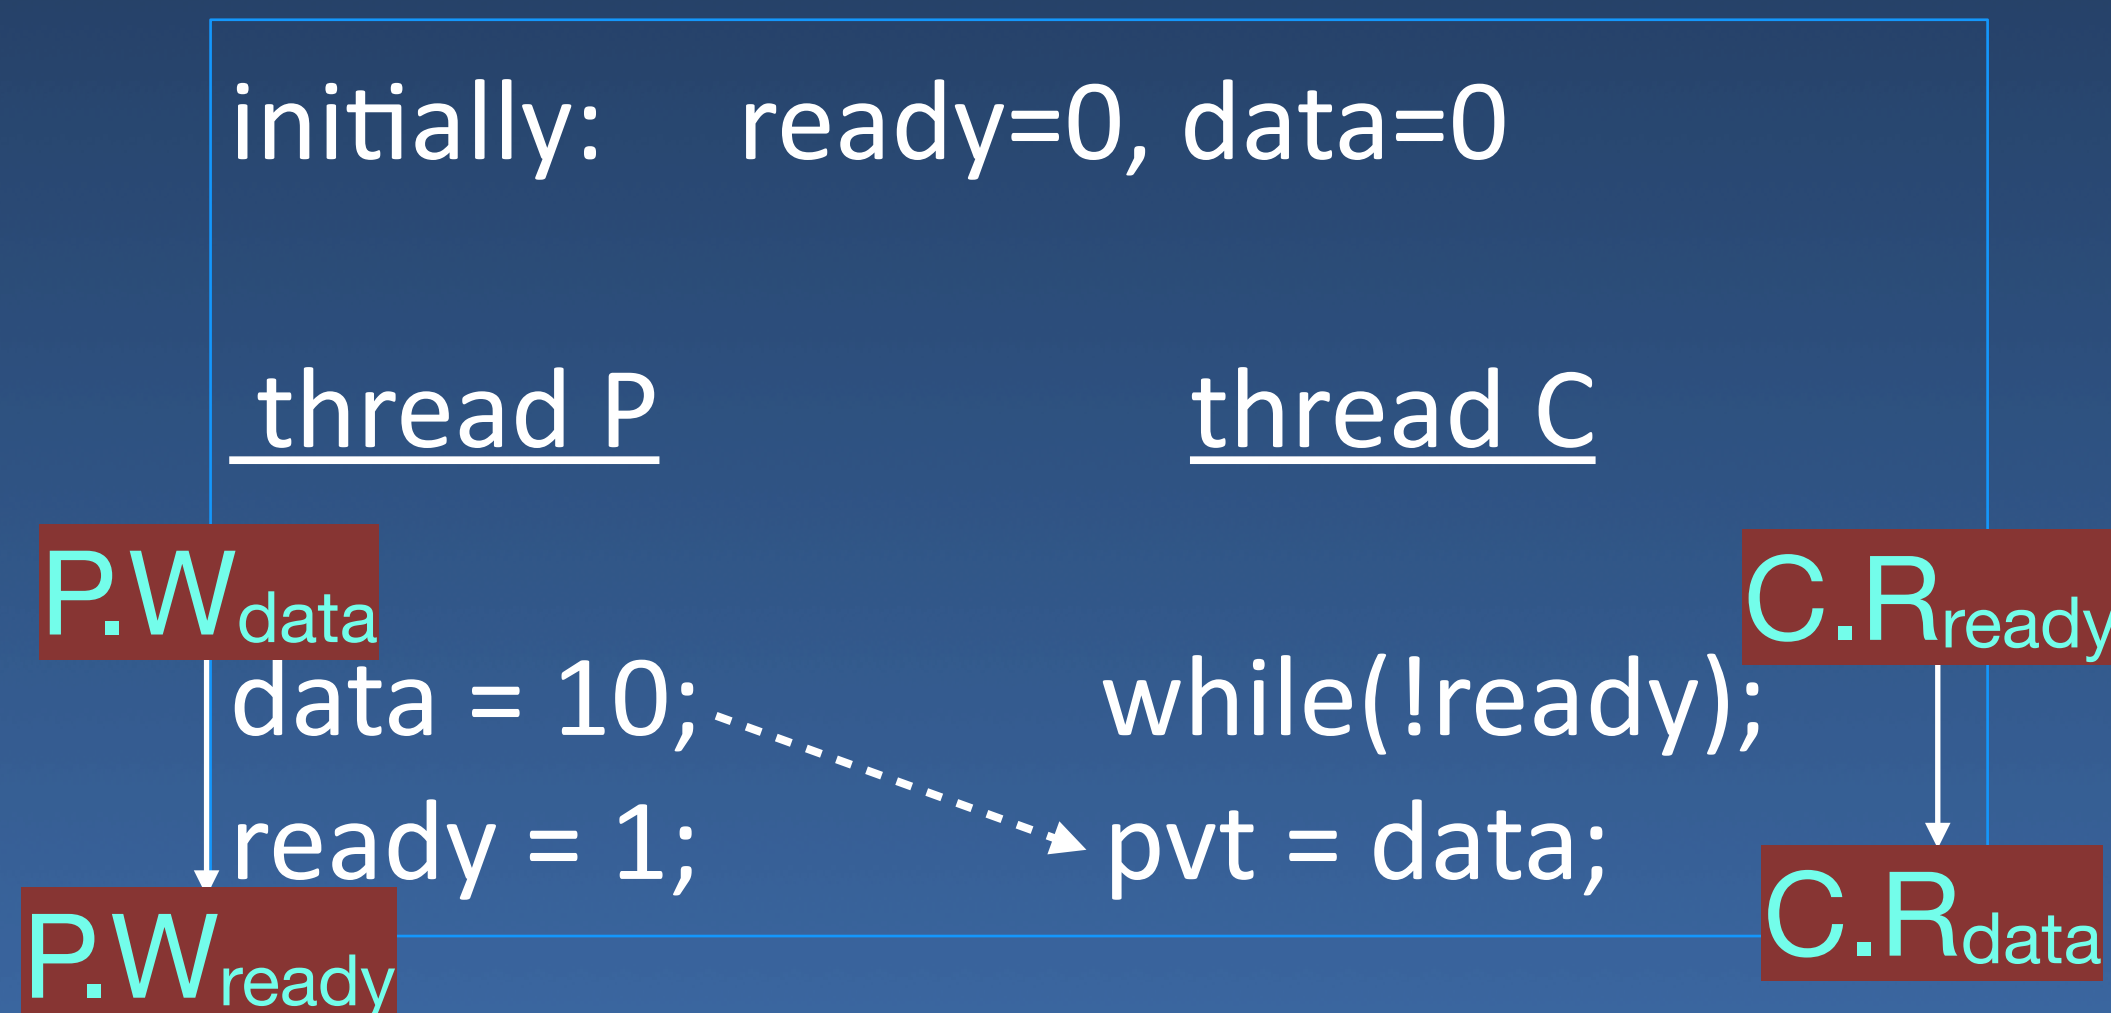
data = 10;  
ready = 1;

while(!ready);  
pvt = data;



## Applying Sequential Consistency

- Threads always see values written by some thread
  - ➔ No garbage (update is atomic)
- The value seen is constrained by thread-order
  - ➔ for every thread



**Show:** If C sees the updated ready (=1),  
C must also see the updated data (=10)



# Applying Sequential Consistency

- Threads always see values written by some thread
  - ➔ No garbage (update is atomic)
- The value seen is constrained by thread-order
  - ➔ for every thread

initially: ready=0, data=0

thread P

data = 10;  
ready = 1;

thread C

while(!ready);  
pvt = data;

P.W<sub>data</sub>

P.W<sub>ready</sub>

C.R<sub>ready</sub>

C.R<sub>data</sub>

If C sees ready =	then C sees data =
0	0 or 10
1	10

**Show:** If C sees the updated ready (=1),  
C must also see the updated data (=10)



# Coherence *vs* Consistency

Initially:  $X = 0$ ;  $Y = 0$ ;

Thread A

$l = Y$

$r = X$

[3]



[0]


Thread B

$X = 5$

$Y \downarrow = 3$


# Coherence *vs* Consistency

Initially:  $X = 0$ ;  $Y = 0$ ;




# Coherence *vs* Consistency

Initially:  $X = 0$ ;  $Y = 0$ ;



# Coherence *vs* Consistency

Initially:  $X = 0$ ;  $Y = 0$ ;





- Consistency is about global state (not per-variable)
  - ➔ Program must not assume higher consistency than available
- Only local memory dependencies visible to compiler/architecture
  - ➔ Can allocate a register or stack entry for some shared variable
  - ➔ Batching of memory transactions
  - ➔ Network can also reorder two memory messages

- Consistency is about global state (not per-variable)
  - Program must not assume higher consistency than available
- Only local memory dependencies visible to compiler/architecture
  - Can allocate a register or stack entry
  - Batching of memory transactions
  - Network can also reorder two memory operations

```
X=1; Y=1; X=2;
```


→ Second write to X may happen before Y's

→ 1st write may never happen

- Consistency is about global state (not per-variable)
  - ➔ Program must not assume higher consistency than available
- Only local memory dependencies visible to compiler/architecture
  - ➔ Can allocate a register or stack entry for some shared variable
  - ➔ Batching of memory transactions
  - ➔ Network can also reorder two memory messages

★ Solutions: Handle inconsistency, force synchronization

## Example: Not SC





## SC Inefficiency

- Hard to implement efficiently
  - Need to enforce serialization of operations
  - No re-ordering of instructions allowed

<u>thread A</u>	<u>thread B</u>	<u>thread C</u>	<u>thread D</u>
<b>x = a</b>	<b>x = b</b>	y1 = x ( <b>b</b> ) y2 = x ( <b>a</b> )	z1=x ( <b>a</b> ) z2=x ( <b>b</b> )

- Some solutions:
  - Allow out-of-order execution, Detect and recover from SC violation
  - Only enforce ordering when required
    - ▶ programmer enforced

# Relaxing SC

initially: ready=0, data=0

thread P


data1 = 1  
data2 = 1  
Say OK

thread C

Wait for OK  
sav1 = data1  
sav2 = data2

## Weak Consistency


- Special **synchronization** accesses are sequentially consistent
- Regular accesses ordered only with respect to synchronization accesses
  - ▶ Before any regular read/write is allowed to be visible to any other thread, all previous synchronization accesses must become visible
  - ▶ Before a synchronization access is allowed to complete, all previous ordinary read/write accesses must be completed
- Suitable for many optimizations





# Weak Consistency

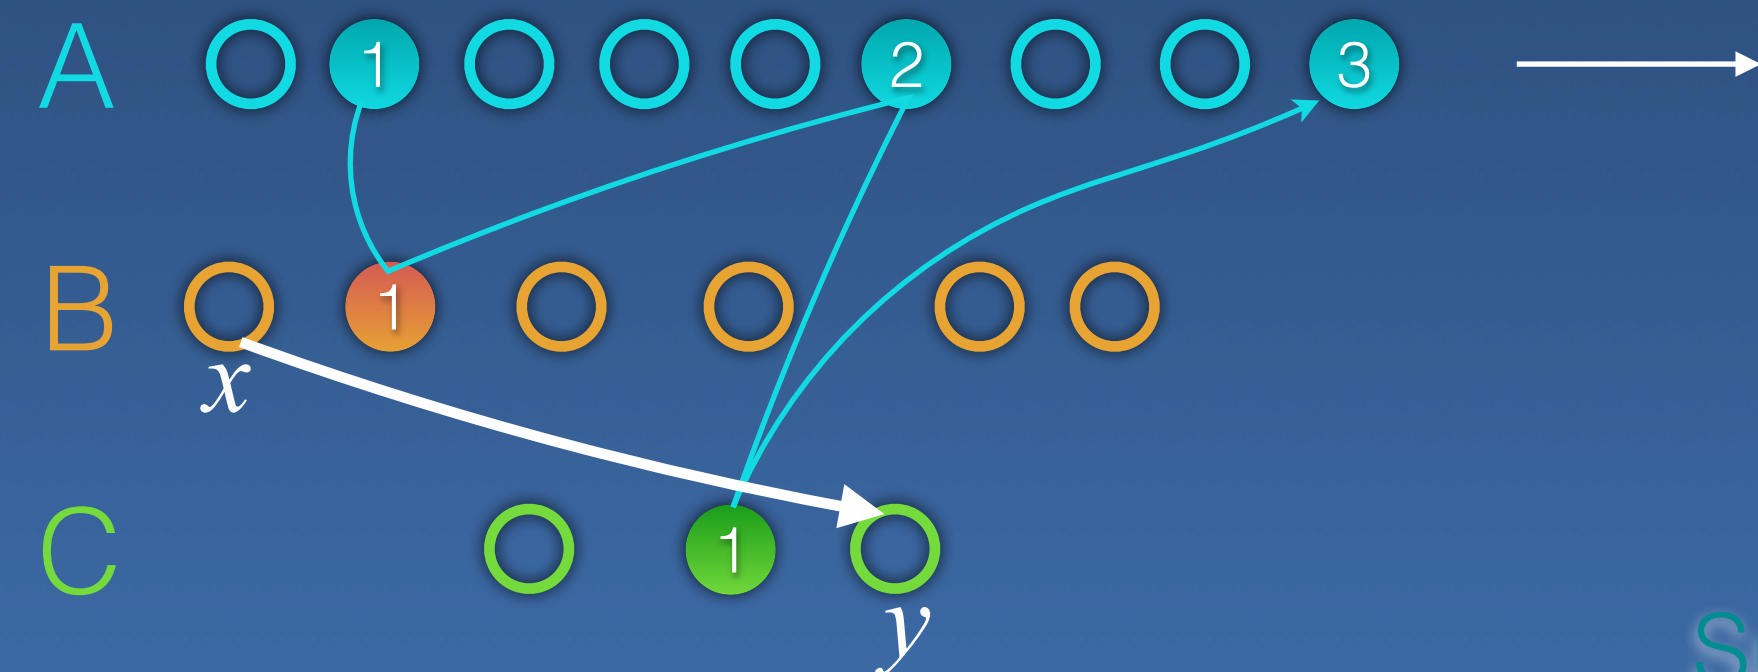
- Special **synchronization** accesses are sequentially consistent
- Regular accesses ordered only with respect to synchronization accesses
  - ▶ Before any regular read/write is allowed to be visible to any other thread, all previous synchronization accesses must become visible
  - ▶ Before a synchronization access is allowed to complete, all previous ordinary read/write accesses must be completed
- Suitable for many optimizations





# Weak Consistency

- Special **synchronization** accesses are sequentially consistent
- Regular accesses ordered only with respect to synchronization accesses
  - ▶ Before any regular read/write is allowed to be visible to any other thread, all previous synchronization accesses must become visible
  - ▶ Before a synchronization access is allowed to complete, all previous ordinary read/write accesses must be completed
- Suitable for many optimizations



# Weak Consistency

- Special **synchronization** accesses are sequentially consistent
- Regular accesses ordered only with respect to synchronization accesses
  - ▶ Before any regular read/write is allowed to be visible to any other thread, all previous synchronization accesses must become visible
  - ▶ Before a synchronization access is allowed to complete, all previous ordinary read/write accesses must be completed
- Suitable for many optimizations

