

COL380

Introduction to  
Parallel & Distributed Programming

```
#pragma omp atomic read/write/update/capture  
X++;
```

- Light-weight critical section
- Only for some basic expressions, e.g,
  - ➔  $x \text{ binop} = \text{expr}$  (no mutual exclusion on expr evaluation)
  - ➔  $X++$
  - ➔  $--X$
  - ➔  $V = X++$

```
#pragma omp atomic read/write/update/capture  
X++;
```

- Light-weight critical section

- Only for some basic operations

→ `x binop= expr` (no `volatile`)

→ `X++`

→ `--X`

→ `V = X++`

## Thread 0

// Produce data

data = 42;

// Set flag to signal Thread 1

#pragma omp atomic write

flag = 1;

## Thread 1

// Busy-wait until flag is signalled

#pragma omp atomic read

myflag = flag

while (myflag != 1) {

#pragma omp atomic read

myflag = flag

}

// Consume data

printf("data=%d\n", data);


`#pragma omp taskwait` `depend(in:x)`

- Wait (suspend) for all children tasks
  - ➔ `depend`  $\Rightarrow$  only wait for some precedent tasks
- Also see:
  - ➔ `#pragma omp taskgroup`




# Critical Section

- A block of code
- Criticality context
  - ➔ wrt other block(s)



```
#pragma omp critical (a_name)
{
    mutually_excluded_code();
}
```

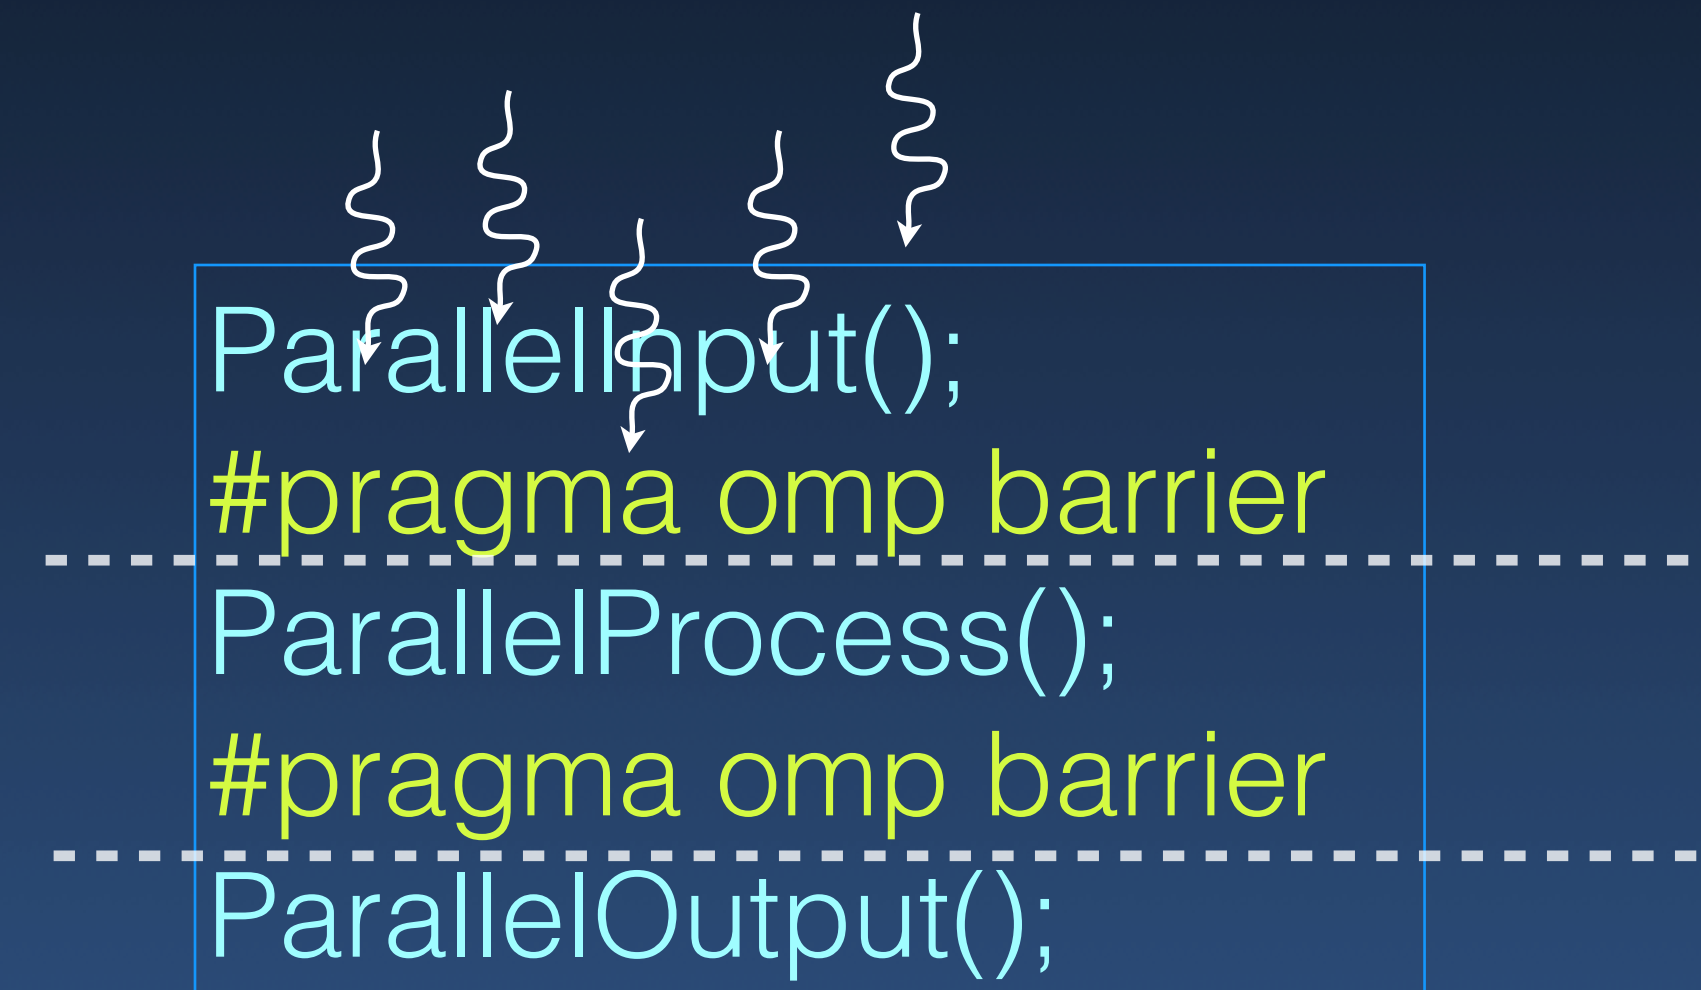
Serialized



```
#pragma omp critical (a_name)
{
    mutually_excluded_code_also();
}
```

# Barrier

- A group of entities
- Wait for all
  - ➔ Any post-barrier computation implies completion of pre-barrier computation in each thread of the group



# Lock

See:

```
int omp_test_lock (omp_lock_t *);
```

```
omp_lock_t lockA;  
omp_init_lock (&lockA);  
...  
omp_destroy_lock (&lockA);
```


- Object: lock
- Actions: Lock and Unlock

Critical Section

```
OperateA(object *A)  
{  
    omp_set_lock(&lockA);  
    Operate_Exclusively(A)  
    omp_unset_lock (&lockA);  
}
```

## Synchronization

- Do Operation  $X$  at time  $T$
- Do we need a precise notion of time to make progress?
  - Always increasing
  - Shared view or individual view?
- Basic synchronization
  - Two (or a set of) events should happen together
  - Any two (from a set of) events should NOT happen together
  - ★ Event  $A$  should happen after event  $B$



Stop and Go



# Causal Ordering

- Define a partial order

- Causality:  $A \rightarrow B \Rightarrow \text{Time}(A) < \text{Time}(B)$     Same as:  $\text{Time}(A) < \text{Time}(B) \Rightarrow B \nrightarrow A$

- Inverse is not required to be true    Means they can be concurrent

Strong causality:  $A \rightarrow B \Rightarrow \text{Time}(A) < \text{Time}(B) \ \&\& \ \text{Time}(A) < \text{Time}(B) \Rightarrow A \rightarrow B$

- Clocks at least must support partial ordering of events

- Can construct total ordering (e.g., by using Process-ID to break tie)

- Possible to build “counters” that can support total order (strong causality)

## Logical-Clock [Lamport's Timestamp algorithm]


- Each entity (process) maintains a counter
  - ➔ increments every 'event,' at its own pace
- Interaction between entities is through messages
  - ➔ Data + counter
- On message receipt:
  - ➔ If recipient counter < received counter
    - ▶ Increase local counter to received counter
    - ▶ Receive is an event, so increment by one

$$A \rightarrow B \Rightarrow \text{Time}(A) < \text{Time}(B)$$

e.g., Lamport Clock

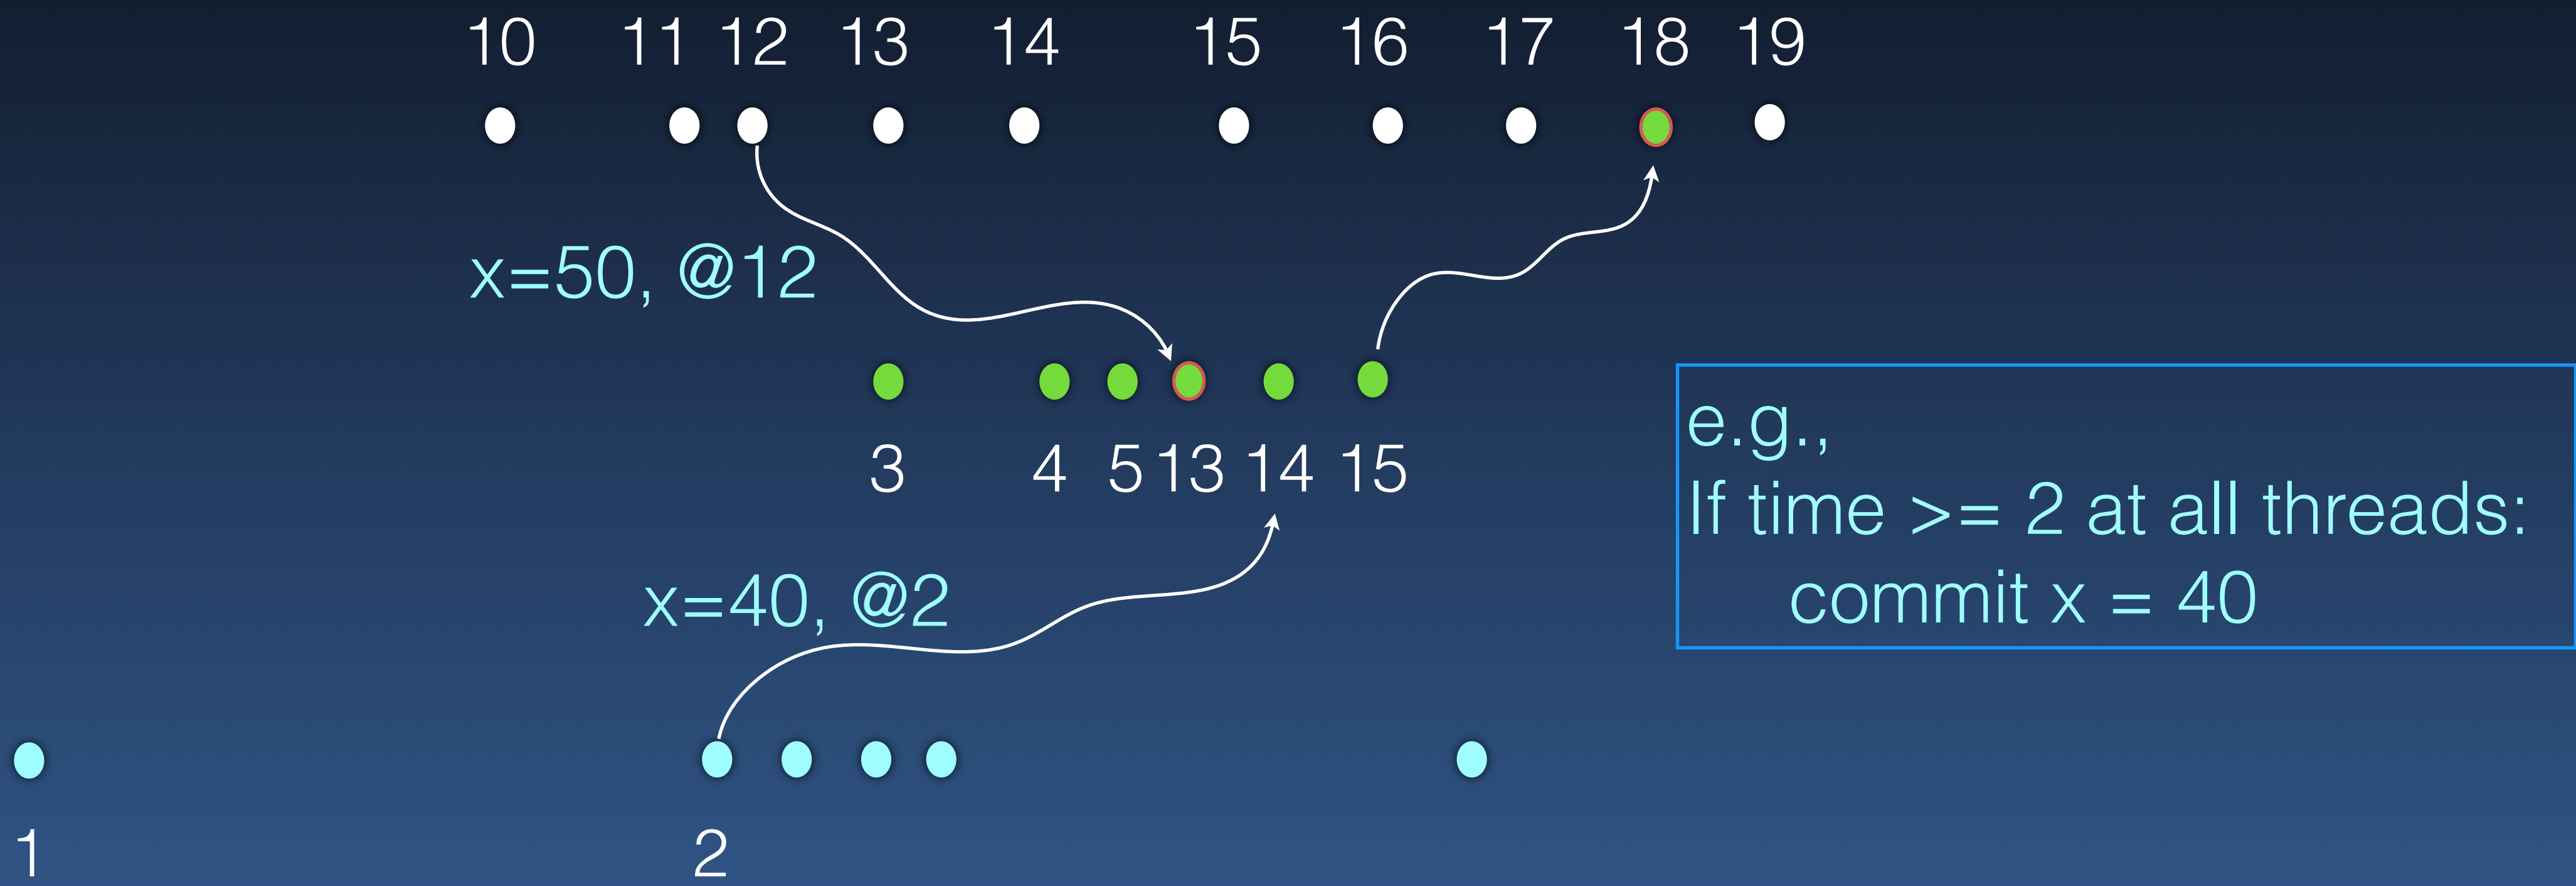


e.g., Lamport Clock





e.g., Lamport Clock



- Memory Fences, consistent memory
  - ➡ Registers
- Atomic operations
  - ➡ Test & Set, Fetch & Add, Compare & Swap
- Critical section, Mutex, Ordered
- Barrier
- Lock
- Wait, Condition variables

# Properties of Synchronization

- Safety, Liveness
- Blocking
- Starvation-free, Deadlock-free, Obstruction-free, Lockfree, Waitfree
- Distributed or Centralized
  - ➔ OS scheduler, Runtime tools

# Synchronization

- Events should happen together

→ Barrier

- Events should NOT happen together

→ Mutual Exclusion, Critical Section

- A should happen before B

→ Conditions, Wait, Ordered

- Lower level Primitives

→ Locks, Semaphores, Registers, Transactional memory

Ease of use? Overhead?

Progress

→ Starvation

→ Deadlock

Blocking vs  
Non-blocking

Busy-wait vs  
OS-scheduled

Fairness

Safety

Liveness



- **Strong Fairness**

- ➔ If any synchronizer is ready infinitely often, it should be executed infinitely often

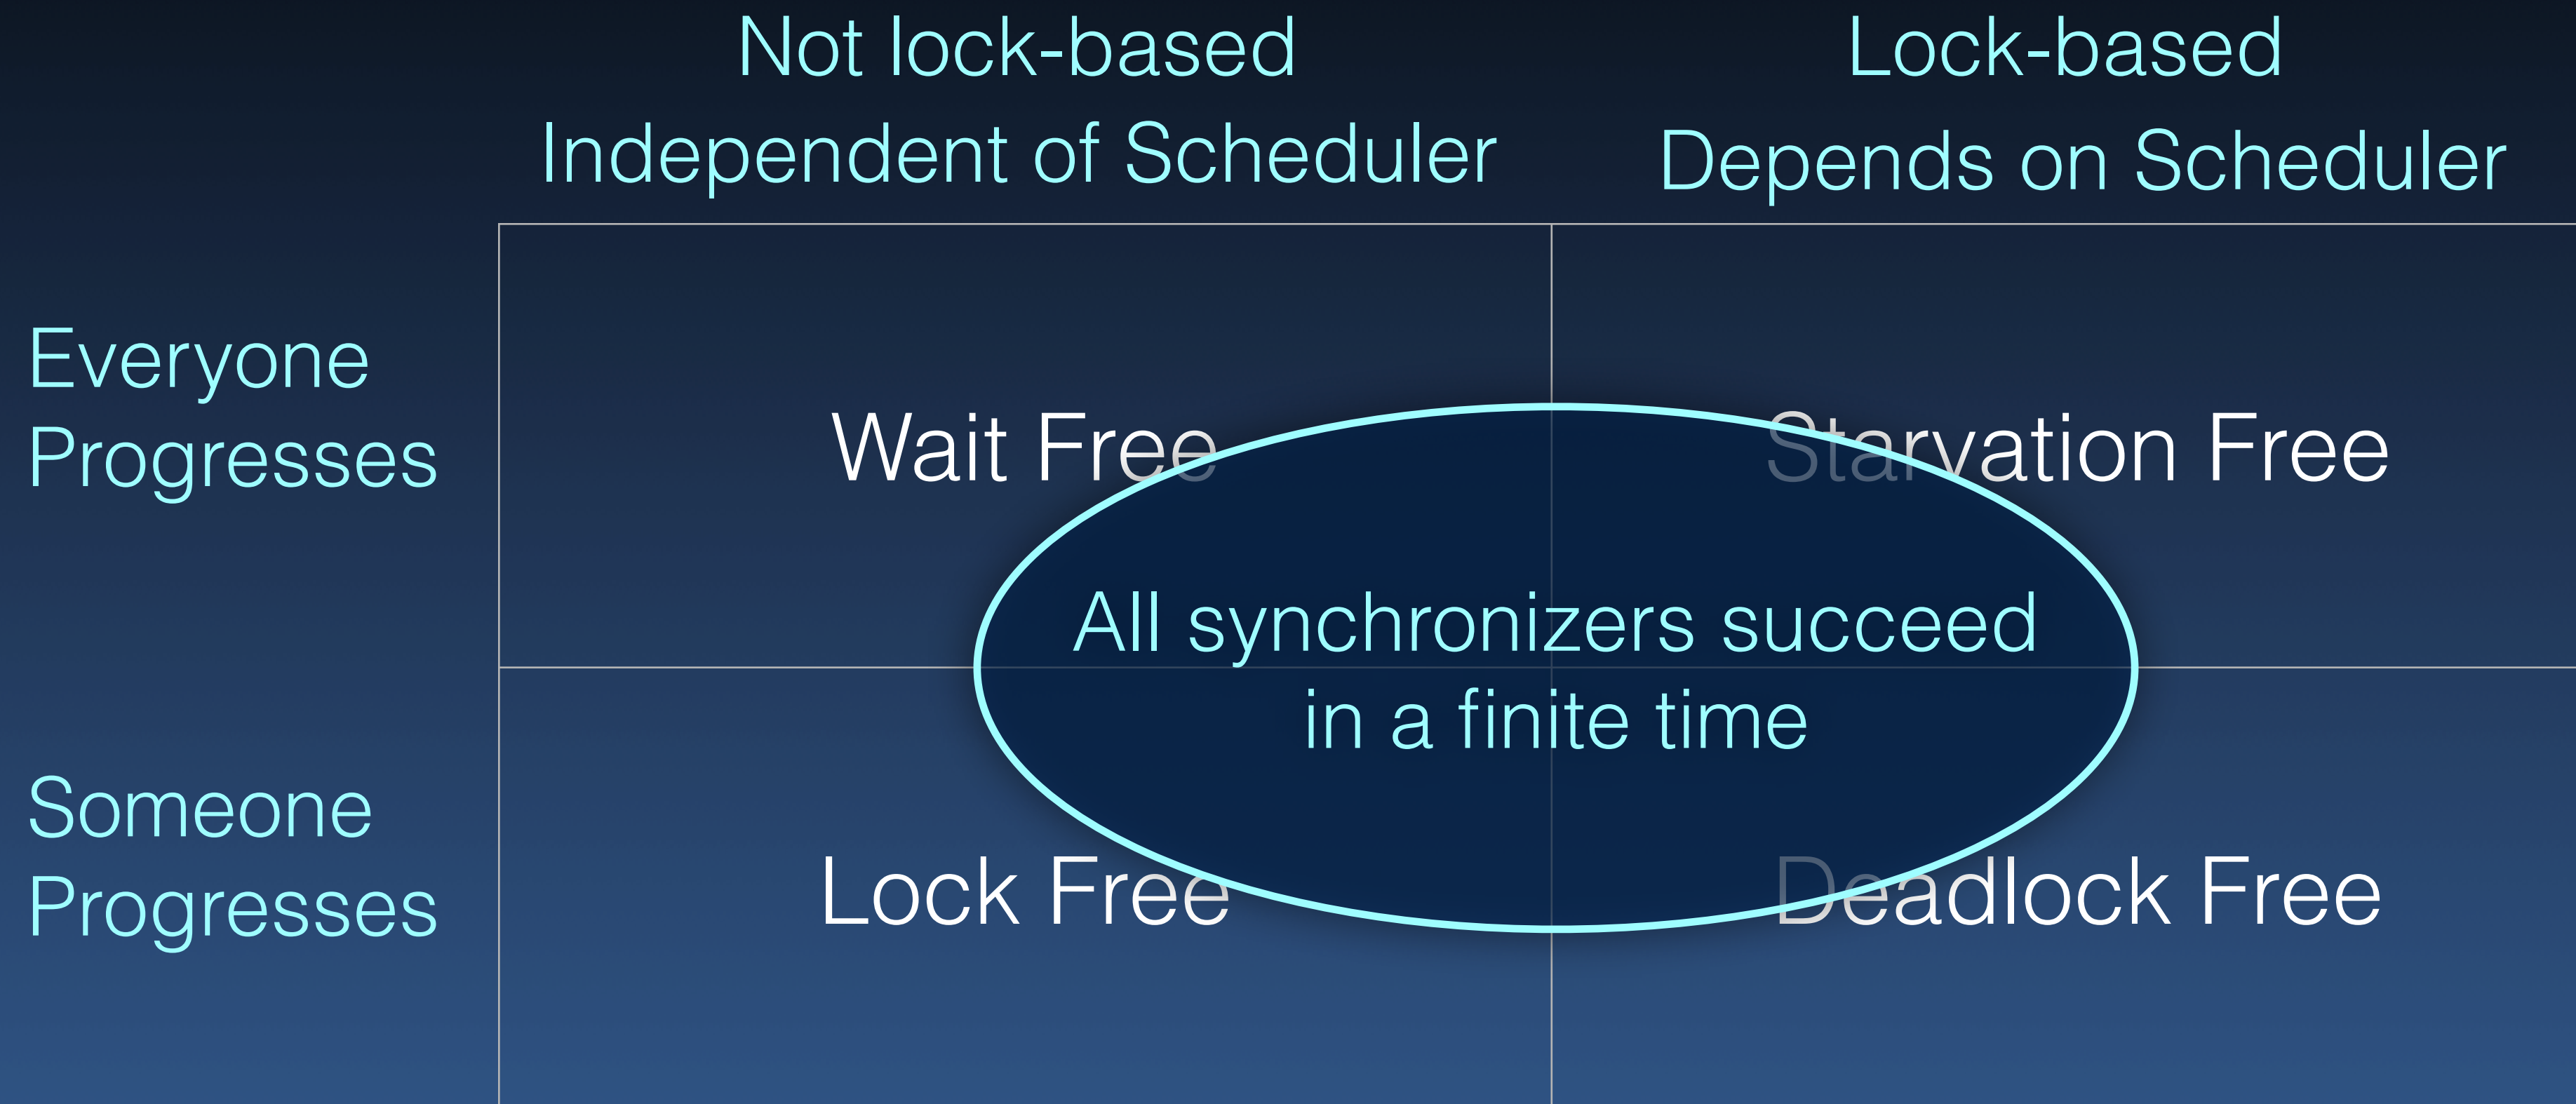
- **Weak Fairness**

- ➔ If any synchronizer is ready, it should be executed eventually

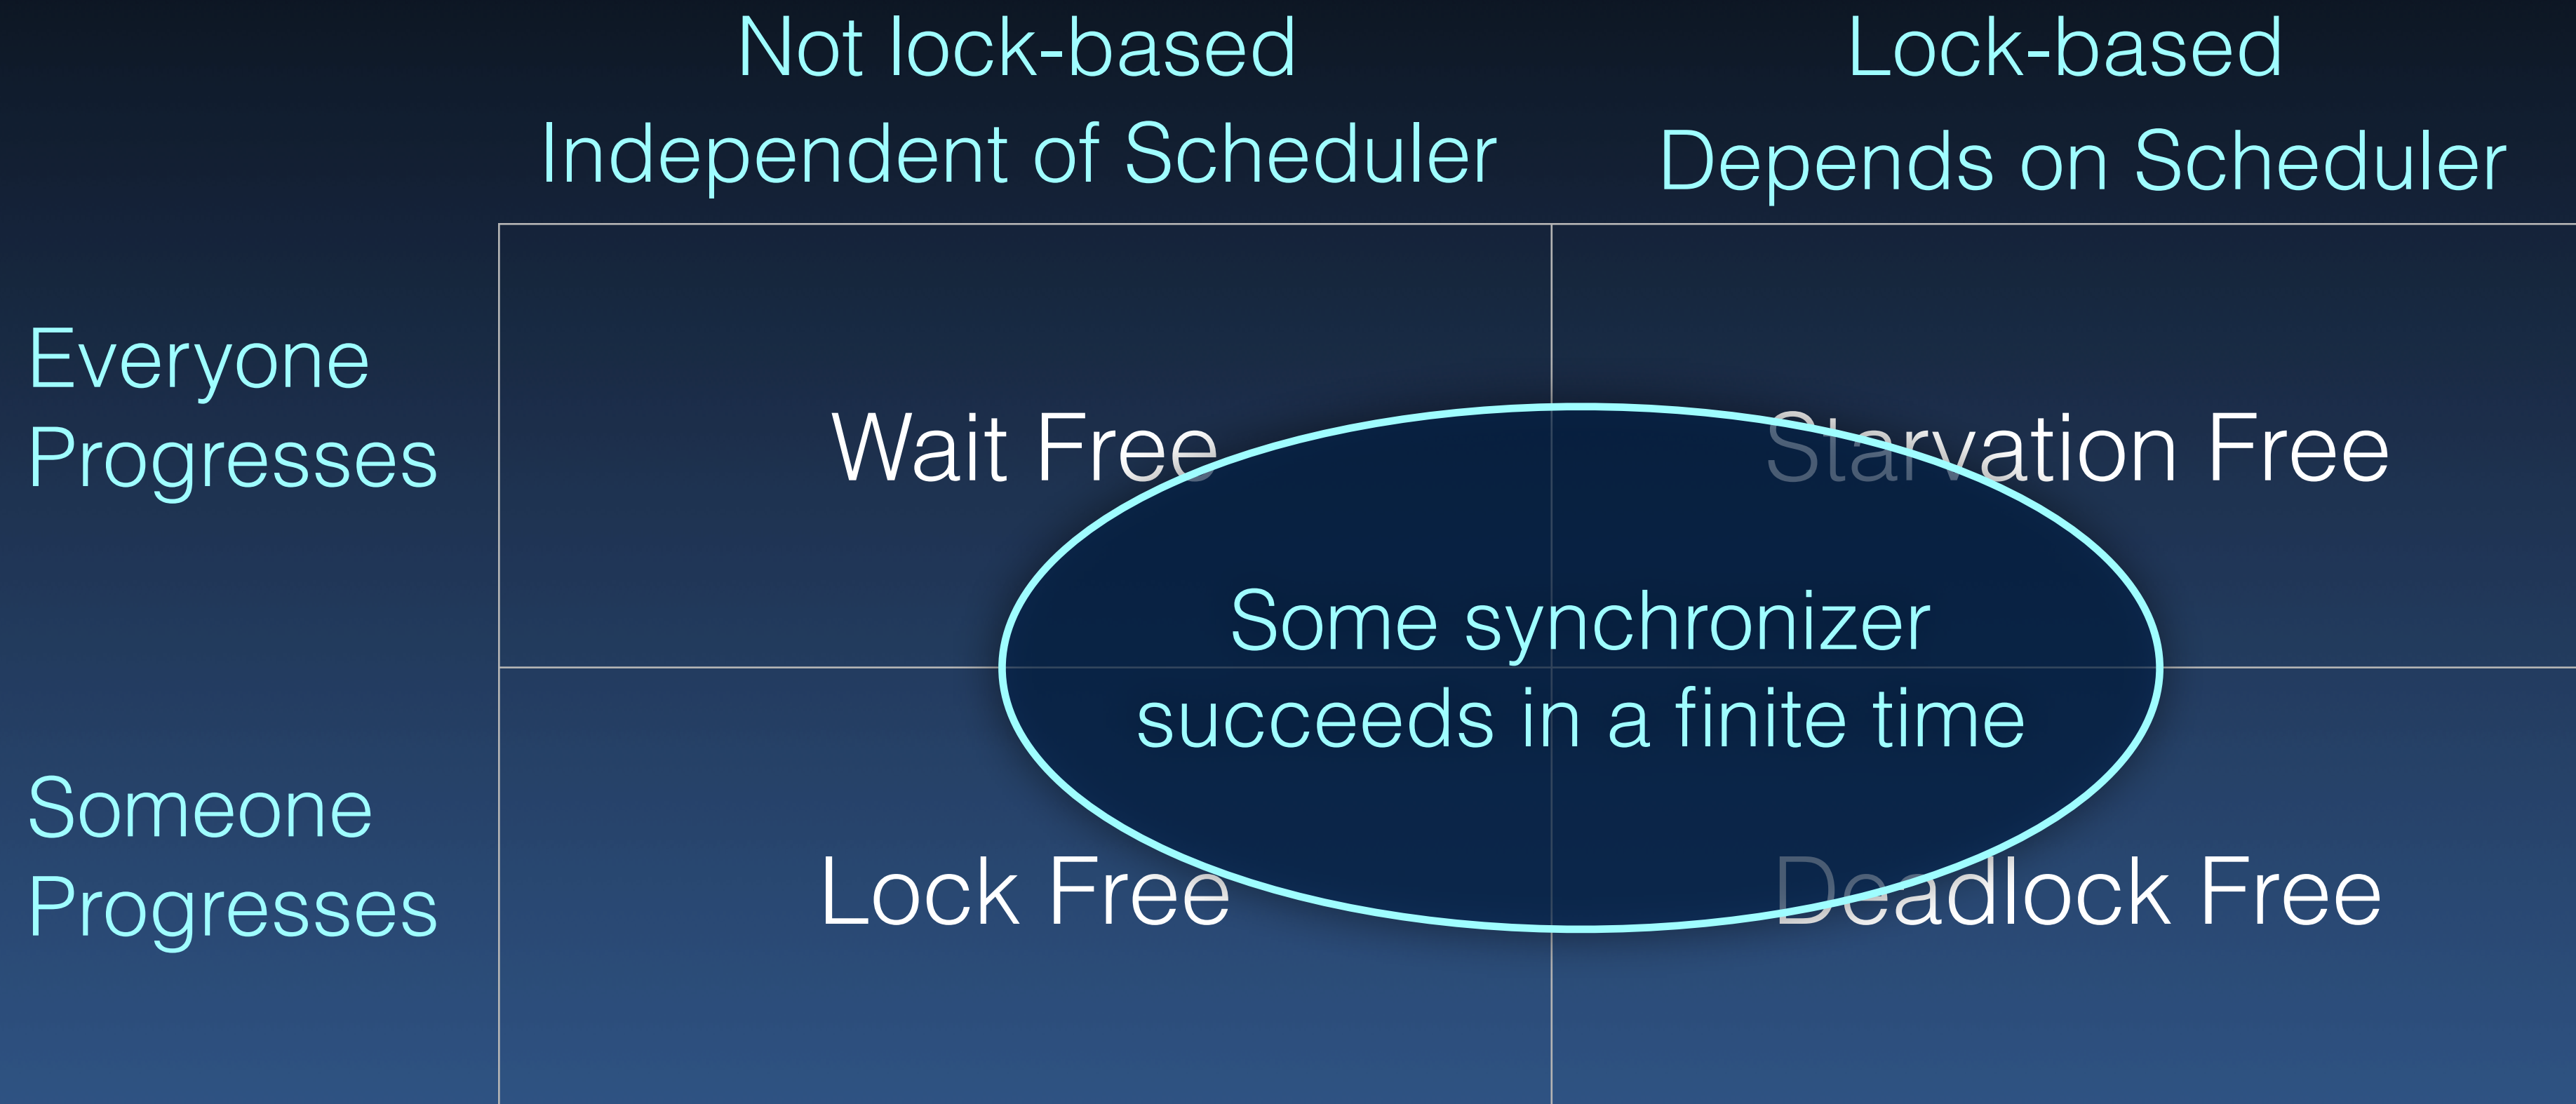
# Progress

	Not lock-based Independent of Scheduler	Lock-based Depends on Scheduler
Everyone Progresses	Wait Free	Starvation Free
Someone Progresses	Lock Free	Deadlock Free

# Progress



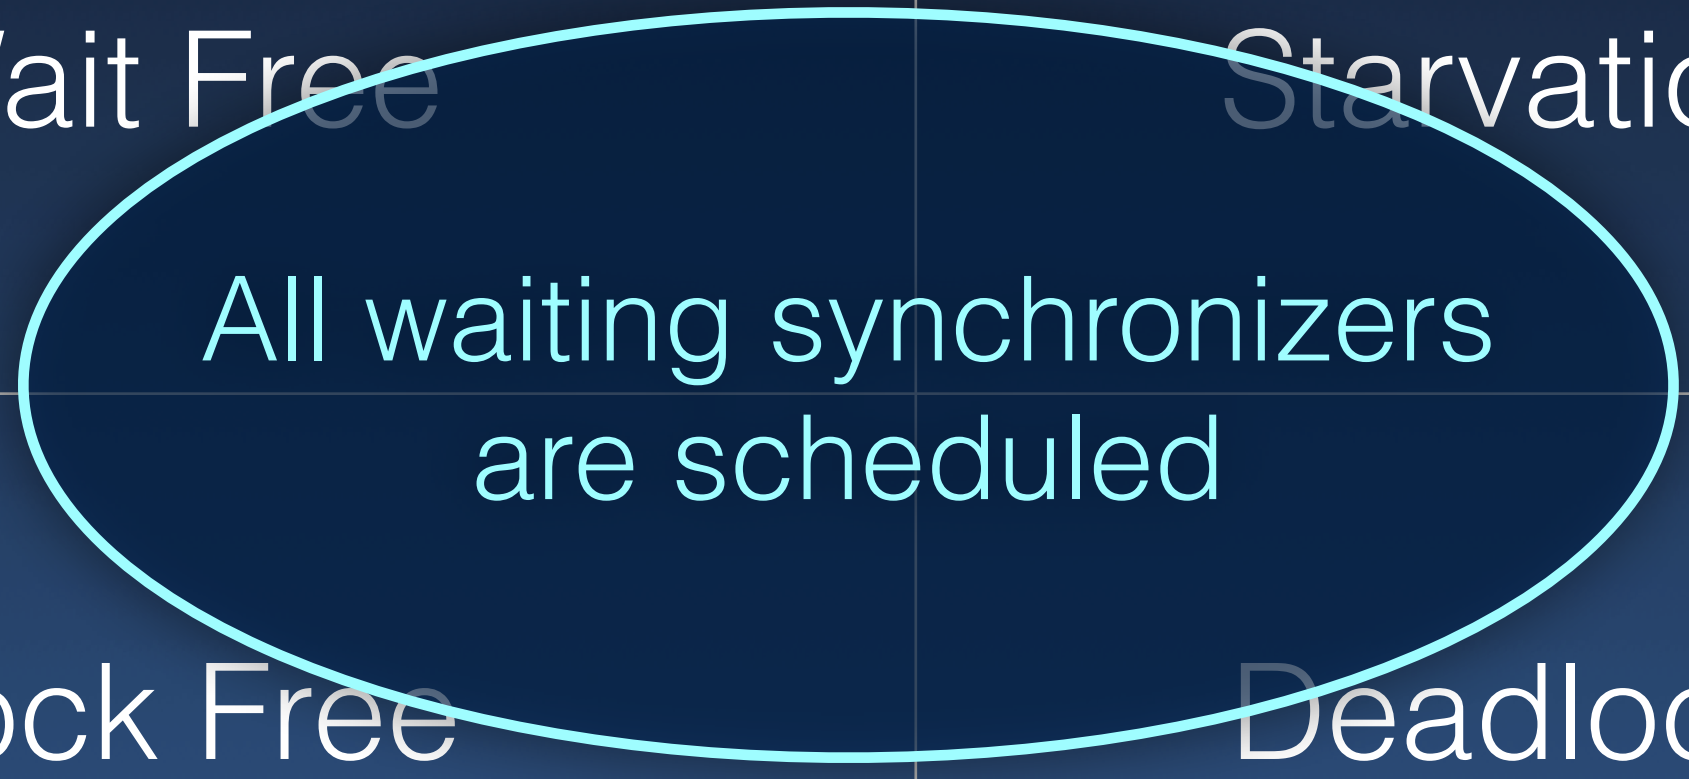
# Progress





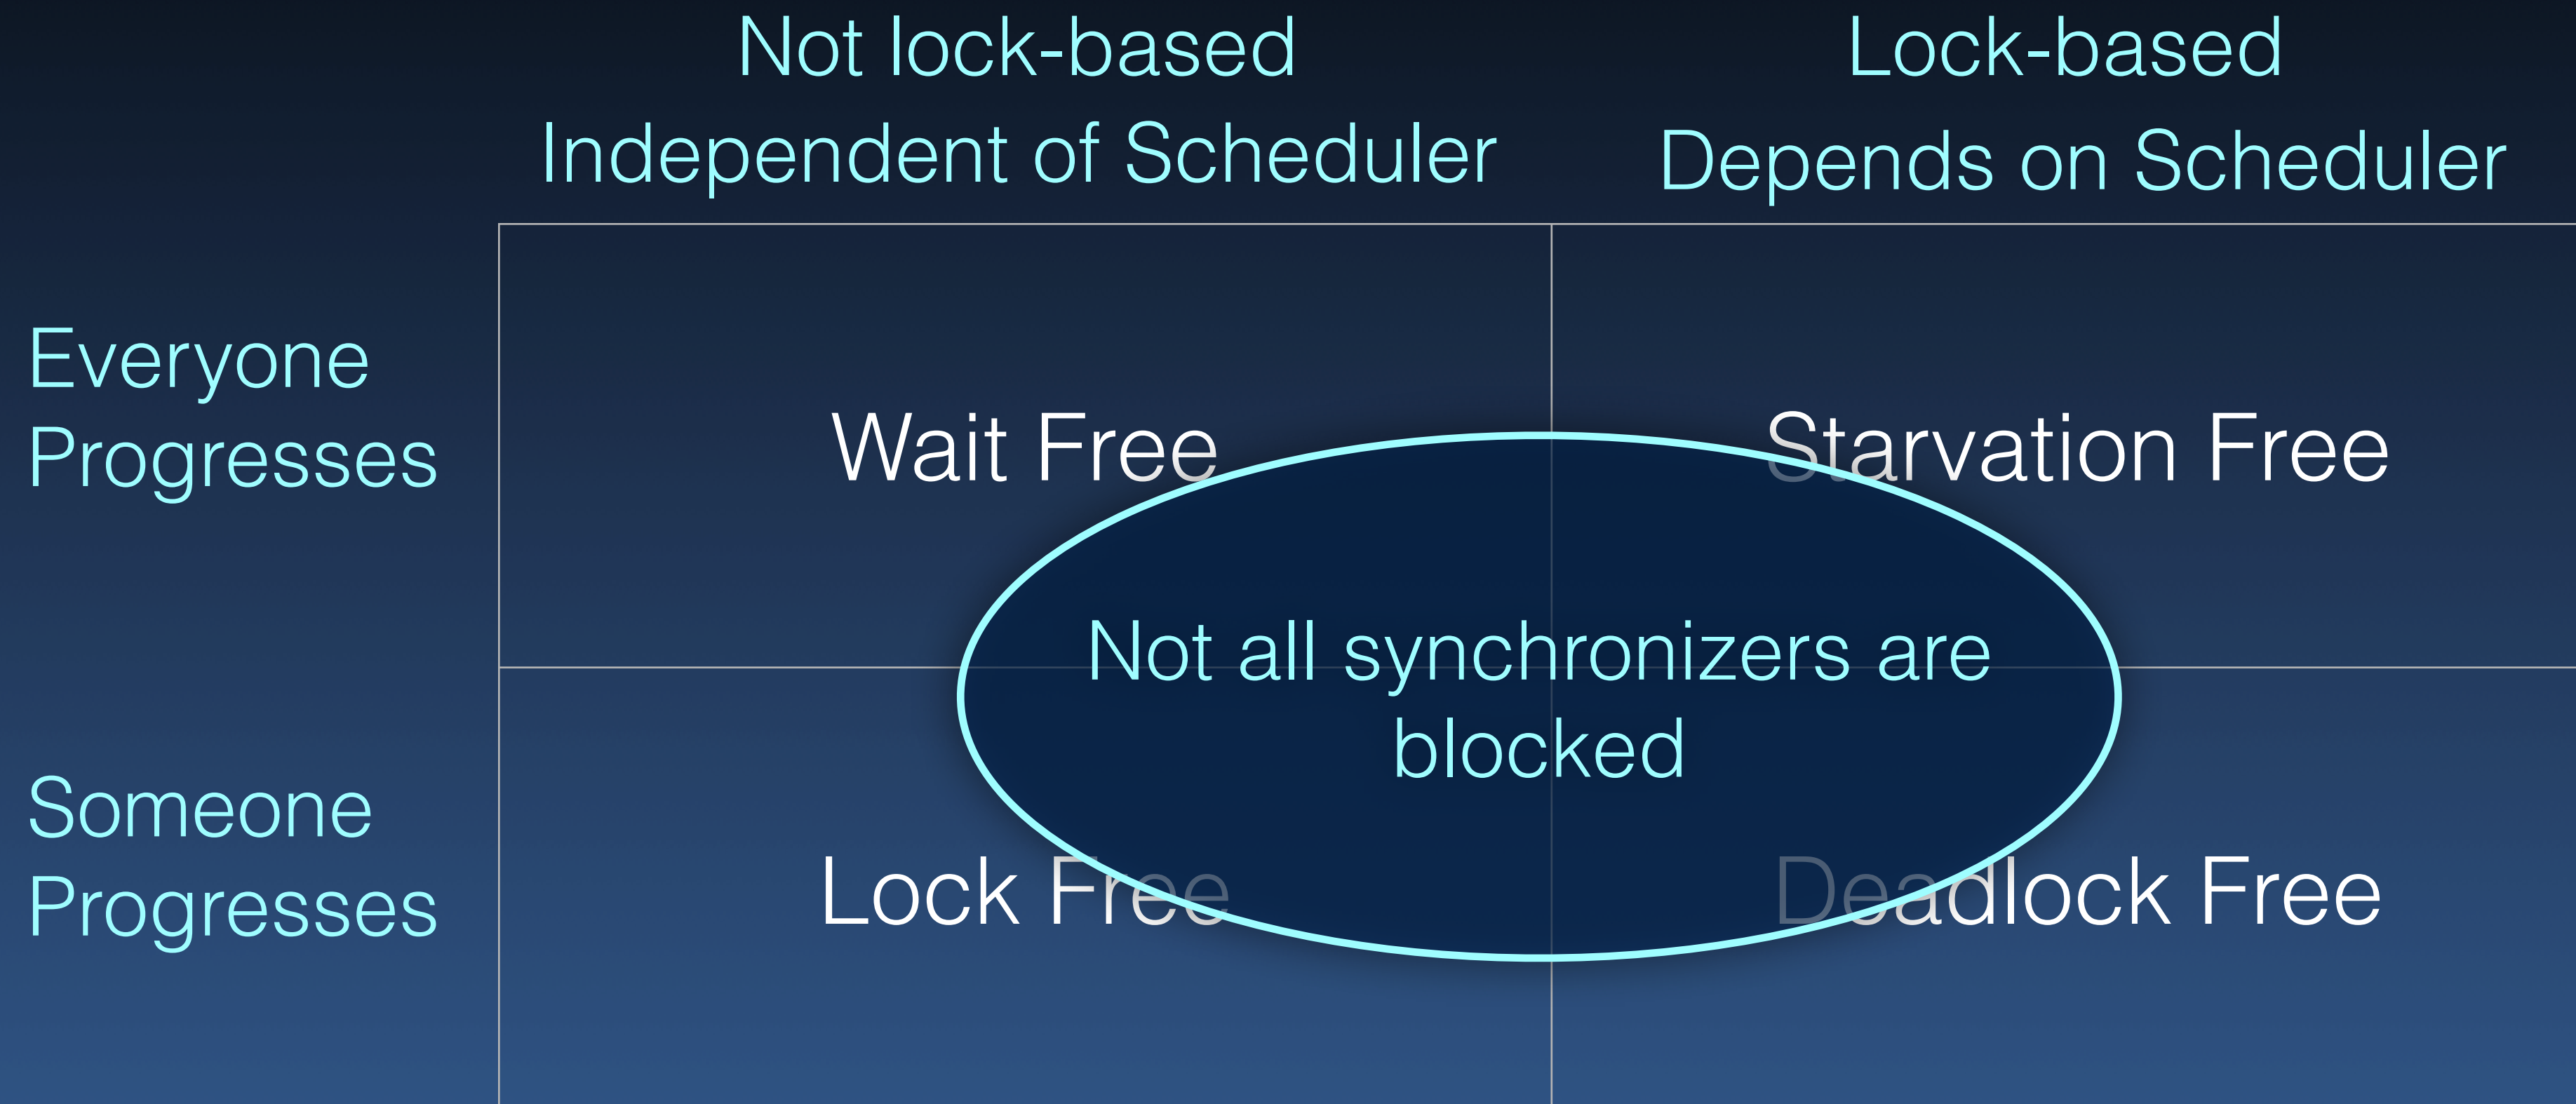
# Progress

	Not lock-based Independent of Scheduler	Lock-based Depends on Scheduler
Everyone Progresses	Wait Free	Starvation Free
Someone Progresses	Lock Free	Deadlock Free



All waiting synchronizers  
are scheduled

# Progress



# Progress

Liveness

Not lock-based Independent of Scheduler	Lock-based Depends on Scheduler
--	------------------------------------

Everyone  
Progresses

Wait Free

Starvation Free

Someone  
Progresses

Lock Free

Deadlock Free

# Scheduler Agnostic Progress

- **Obstruction-free**

- ➔ A given tasks will complete if no other tasks are running (interfering)

- **Lockfree**

- ➔ Some task will complete (no matter what others do)

Complete in a finite number of its own steps

- **Waitfree**

- ➔ Each task will complete (no matter what others do)

Assume others do not behave maliciously