


COL380

Introduction to  
Parallel & Distributed Programming

- Memory Fences, consistent memory
  - ➡ Registers
- Atomic operations
  - ➡ Test & Set, Fetch & Add, Compare & Swap
- Critical section, Mutex, Ordered
- Barrier
- Lock
- Wait, Condition variables

# Compare & Exchange (aka Compare & Swap)



```
std::atomic<int> var(0);
```

```
var.compare_exchange_strong(expected, newval);
```

```
// Atomically:  
// t = var.load();  
// if(t == expected) {  
//     var.store(newval);  
//     return true  
// } else {  
//     return false  
// }
```

```
#pragma atomic  
var++;
```

```
#pragma omp atomic capture compare  
{  
    old = svar;  
    if (old == expected) svar = newval;  
}  
// old == expected ⇒ success
```

# Compare & Exchange (aka Compare & Swap)


```
std::atomic<int> var(0);
```

```
var.compare_exchange_strong(expected, newval);
```

```
// Atomically:  
// t = var.load();  
// if(t == expected) {  
//     var.store(newval);  
//     return true  
// } else {  
//     return false  
// }
```

```
#pragma atomic  
var++;
```

```
#pragma omp atomic capture compare  
{  
    old = svar;  
    if (old == expected) svar = newval;  
}  
  
// old == expected ⇒ success
```





## Compare & Exchange (aka Compare & Swap)


```
std::atomic<int> var(0);
```

```
var.compare_exchange_strong(expected, newval);
```

```
// Atomically:  
// t = var.load();  
// if(t == expected) {  
//     var.store(newval);  
//     return true  
// } else {  
//     return false  
// }
```

```
#pragma atomic  
var++;
```

```
#pragma omp atomic capture compare  
{  
    old = svar;  
    if (old == expected) svar = newval;  
}  
  
// old == expected ⇒ success
```




```
std::atomic<int> var(0);
```

```
var.compare_exchange_strong(expected, newval);
```

```
// Atomically:  
// t = var.load();  
// if(t == expected) {  
//     var.store(newval);  
//     return true  
// } else {  
//     return false  
// }
```

```
#pragma atomic  
var++;
```

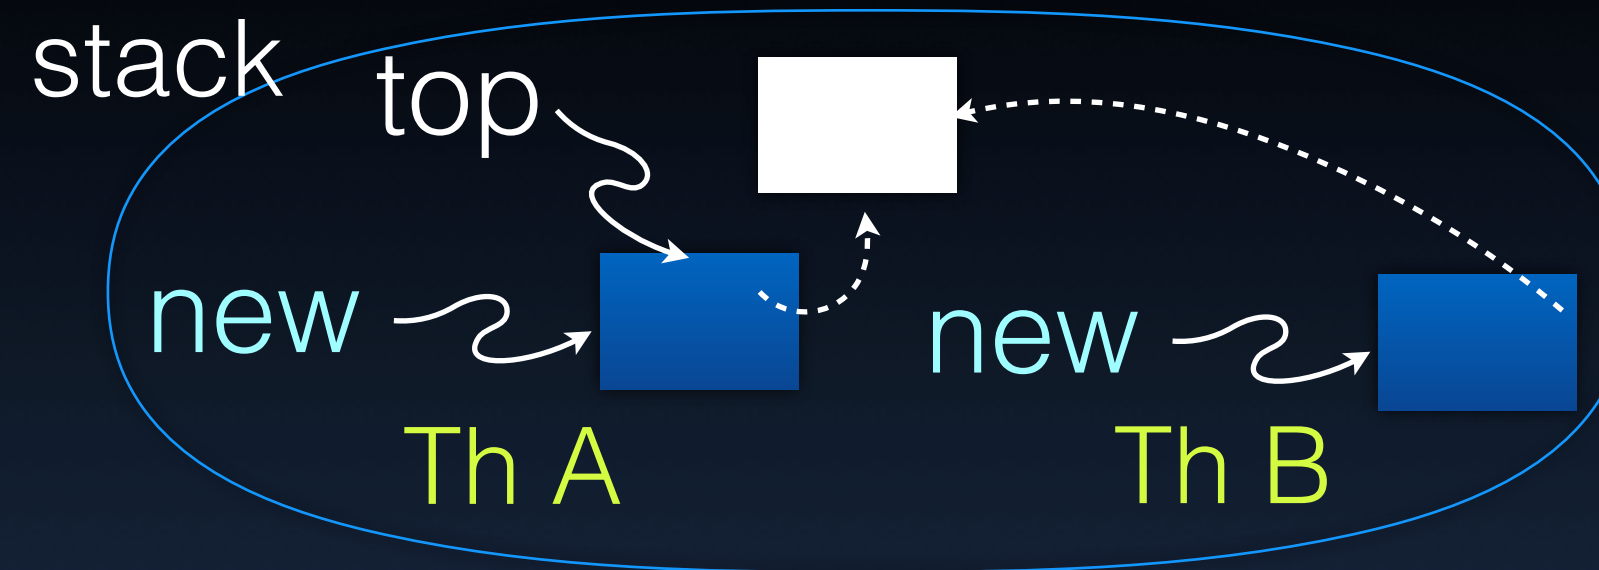
```
#pragma omp atomic capture compare  
{  
    old = svar;  
    if (old == expected) svar = newval;  
}  
  
// old == expected ⇒ success
```



Compare & Exchange  
(aka Compare & Swap)

## Compare & Exchange (aka Compare & Swap)

```
std::atomic<int> var(0);
```



```
var.compare_exchange_strong(expected, newval);
```

```
// Atomically:
```

```
#pragma atomic
```

```
// std::atomic<node<T>*> top;  
// ...  
// void push(const T& data) {  
//     node<T>* new_node = new node<T>(data);  
//     // put the current value of top into new_node->next  
//     new_node->next = top.load();  
//     // Update top to point to the new node  
//     top.store(new_node);  
// }
```


```
are
```

```
l;
```



## Compare & Exchange (aka Compare & Swap)

```
std::atomic<int> var(0);
```



```
var.compare_exchange_strong(expected, newval);
```

```
// Atomically:
```

```
#pragma atomic
```

```
// std::atomic<node<T>*> top;  
// ...  
// void push(const T& data) {  
//     node<T>* new_node = new node<T>(data);  
//     // put the current value of top into new_node->next  
//     new_node->next = top.load();  
//     // Update top to point to the new node  
//     top.store(new_node);  
// }
```


```
are
```

```
l;
```



# Compare & Exchange

```
std::atomic<int> var(0);
```



```
var.compare_exchange_strong(expected, newval);
```

```
// Atomically:
```

```
#pragma atomic
```

```
// std::atomic<node<T>*> top;
```

```
// ...
```

```
// ...
```

```
// void push(const T& data) {
```

```
//     node<T>* new_node = new node<T>(data);
```

```
//
```

```
//     // put the current value of top into new_node->next
```

```
//     do new_node->next = top.load();
```

```
//
```

```
//     // make new_node the top, as long as top still equals new_node->next
```

```
//     while(!top.compare_exchange_strong(new_node->next, new_node));
```


```
}
```

```
are
```

```
l;
```

# Compare & Exchange

```
std::atomic<int> var(0);
```



```
var.compare_exchange_strong(expected, newval);
```

```
// Atomically:
```

```
#pragma atomic
```

```
// std::atomic<node<T>*> top;
```

```
//
```

```
// ...
```

```
// void push(const T& data) {
```

```
//     node<T>* new_node = new node<T>(data);
```

```
//
```

```
//     // put the current value of top into new_node->next
```

```
//     do new_node->next = top.load();
```

```
//
```

```
//     // make new_node the top, as long as top still equals new_node->next
```

```
//     while(!top.compare_exchange_strong(new_node->next, new_node));
```

```
}
```

are

l;

## Fetch & Add

```
#pragma omp atomic capture
{
    old = svar;
    svar += tval;
}
```

```
#pragma omp atomic capture
{
    old = slock;
    slock += 1;
}
if(old == 0) {
    criticalSection();
    #pragma omp atomic
    slock--;
} else {
    havefuninthesun();
}
```



## Test & Set

```
#pragma omp atomic capture
{
    old = svar;
    svar = tval;
}
```

```
#pragma omp atomic capture
{
    old = slock;
    slock = 1;
}
if(old == 0) {
    criticalSection();
    #pragma omp atomic write
    slock=0;
} else {
    havefuninthesun();
}
```

# Condition Variable

- Raise the condition
- Wait for a condition to 'hold'

```
Produce();  
acv.notify_one();
```

```
std::condition_variable acv;  
...  
  
std::unique_lock<std::mutex> alock(amutex);  
acv.wait(alock);  
.. Condition Holds Now ..  
Consume();
```

# Condition Variable

```
void producer(std::condition_variable *cv) {  
    while(1) {  
        produce();  
        cv->notify_one();  
    }  
}
```

```
void consumer(std::mutex *mtx,  
              std::condition_variable *cv) {  
    while(1) {  
        std::unique_lock<mutex> lock(*mtx);  
        cv->wait(lock, 1);  
        consume();  
    }  
}
```

```
{  
    std::mutex mtx;  
    std::condition_variable cv;  
  
    thread p(producer, &cv);  
    thread c(consumer, &mtx, &cv);  
  
    c.join(); p.join();  
}
```



# Condition Variable

```
void producer(std::condition_variable *cv) {  
    while(1) {  
        produce(); counter++;  
        cv->notify_one();  
    }  
}
```


```
void consumer(std::mutex *mtx,  
              std::condition_variable *cv) {  
    while(1) {  
        std::unique_lock<mutex> lock(*mtx);  
        cv->wait(lock, [] { return counter > 0; }  
        consume();  
        counter--;  
    }  
}
```

```
std::atomic<int> counter{0};
```

```
{  
    std::mutex mtx;  
    std::condition_variable cv;  
  
    thread p(producer, &cv);  
    thread c(consumer, &mtx, &cv);  
  
    c.join(); p.join();  
}
```

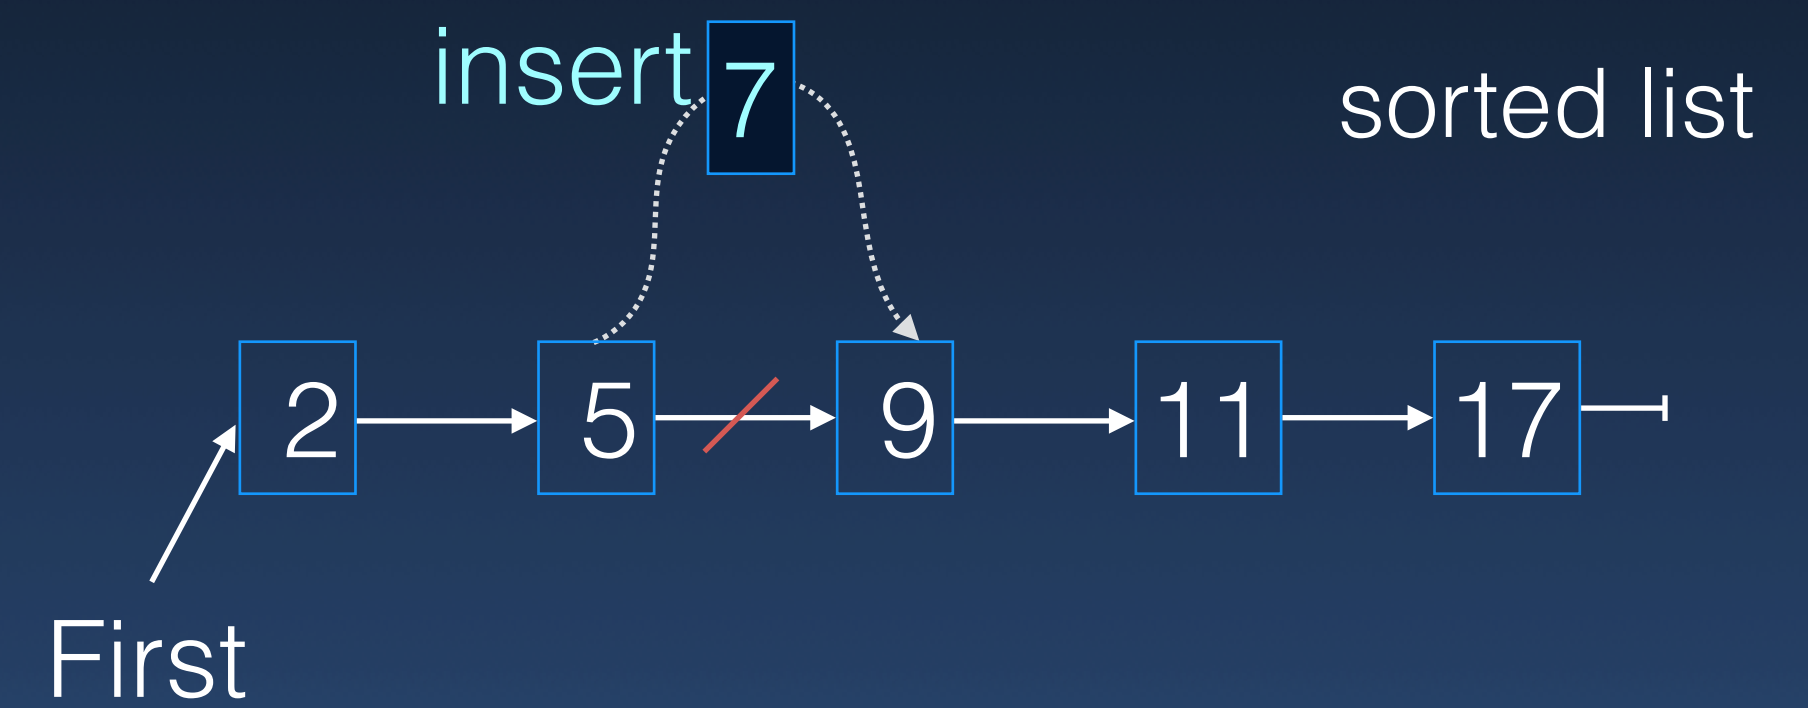
# Lock

- Lock “resources”
- Process
- Unlock “resources”



# Lock


- Lock “resources”
- Process
- Unlock “resources”





# Lock


Correctness?  
“Sequential Equivalence”



- Lock “resources”
- Process
- Unlock “resources”


# Lock

- Lock “resources”
- Process
- Unlock “resources”



# Lock


- Lock “resources”
- Process
- Unlock “resources”





# Lock

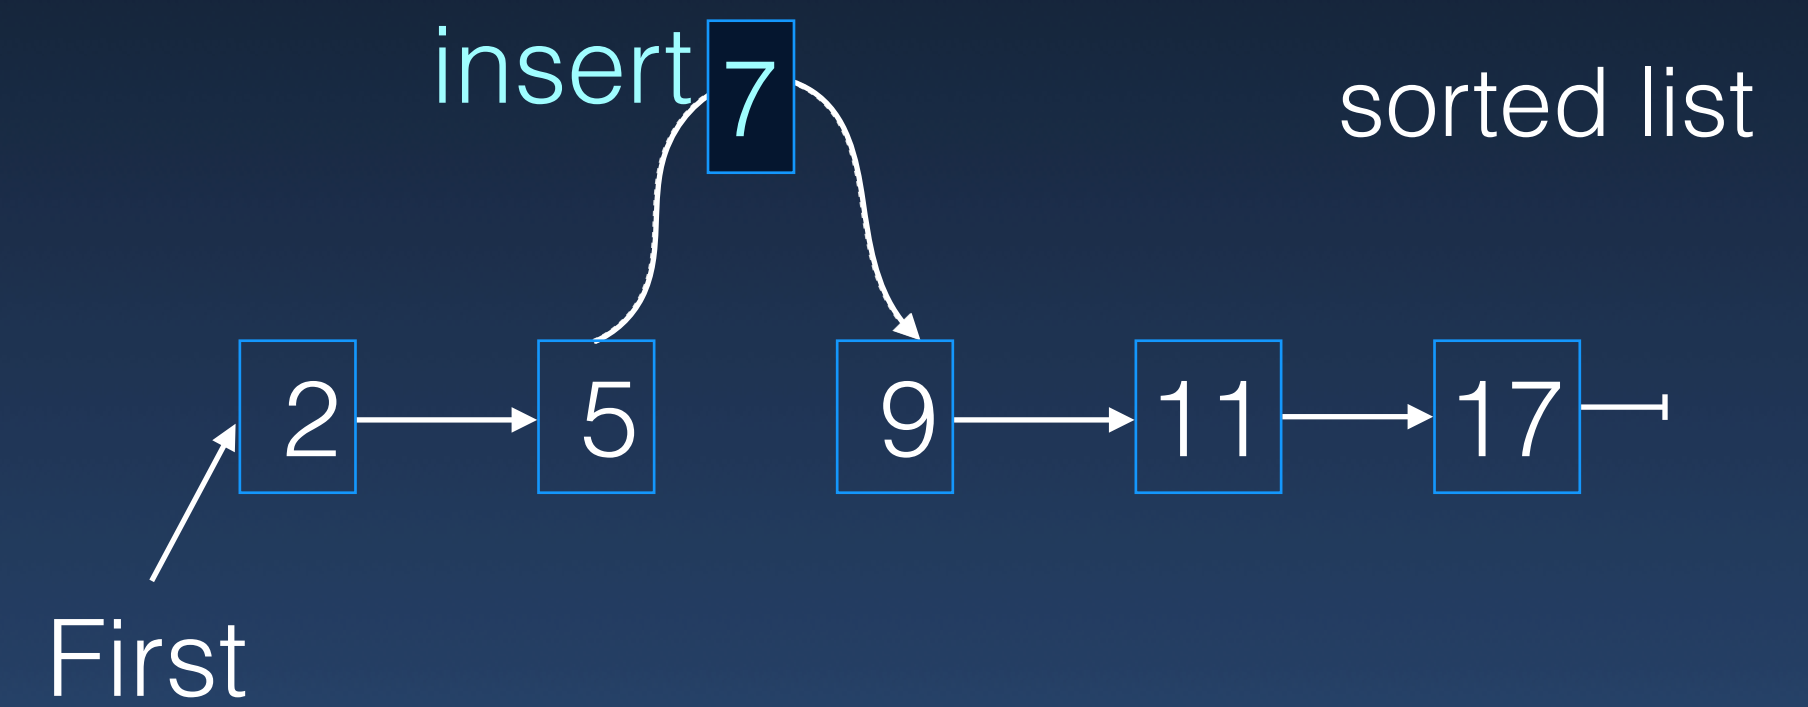
- Lock “resources”
- Process
- Unlock “resources”



# Lock

- Lock “resources”
- Process
- Unlock “resources”

lock(lockA)



# Lock

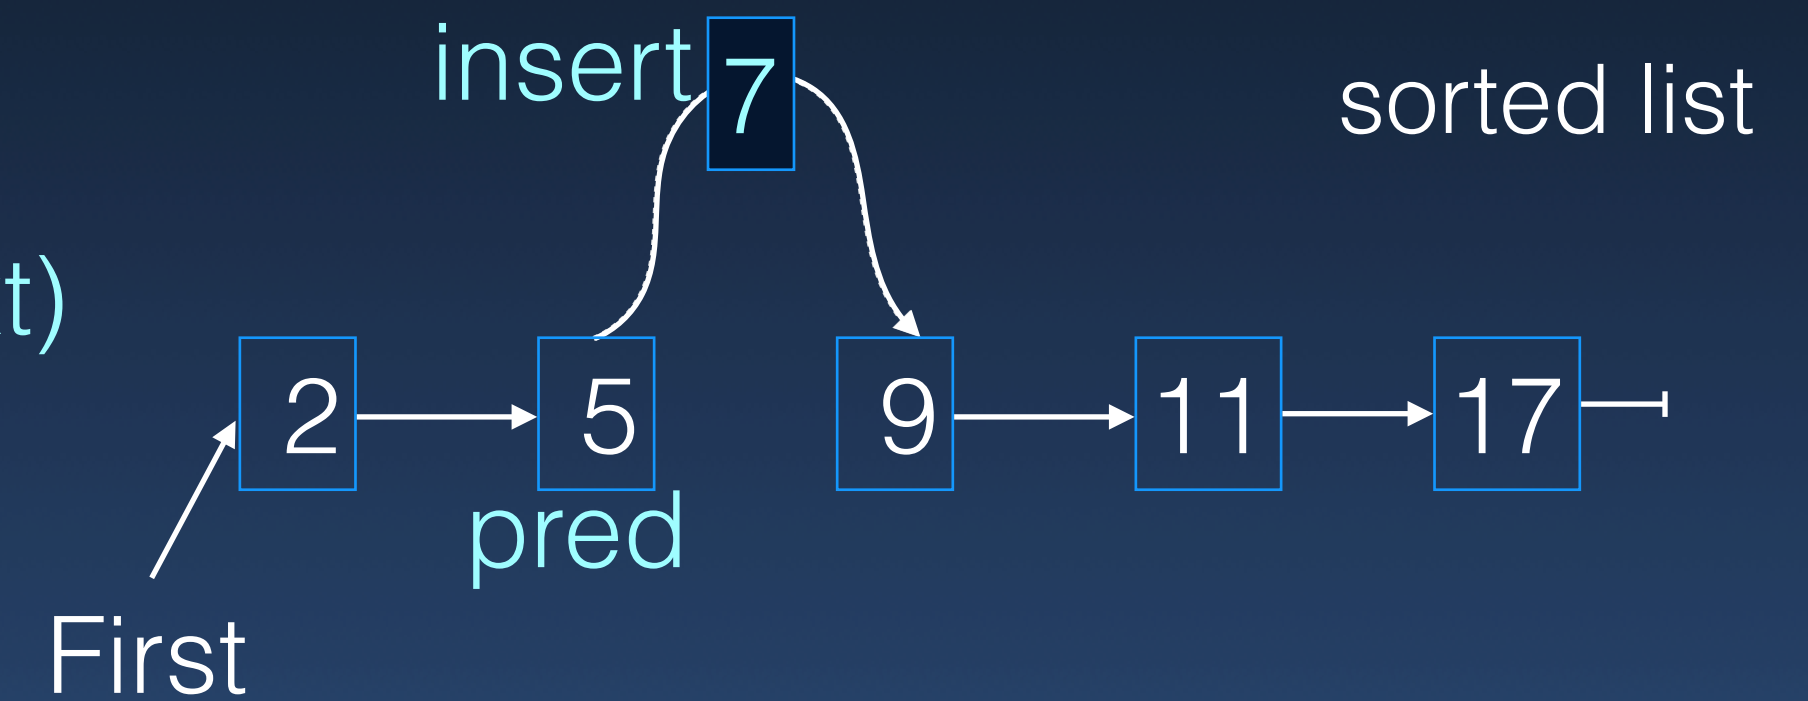
- Lock “resources”
- Process
- Unlock “resources”

`lock(lockA)`

`pred = Find(key)`

`pred.nxt = Node(key, pred.nxt)`

`unlock(lockA)`



# Lock

- Lock “resources”
- Process
- Unlock “resources”

<Request> [?block] <Acquired>


lock(*lockA*)

pred = Find(key)

pred.nxt = Node(key, pred.nxt)

unlock(*lockA*)

<Release> [schedule]





# Lock

- Lock “resources”

<Request> [?block] <Acquired>

lock(lockA)

- Process


pred = Find(key)

pred.nxt = Node(key, pred.nxt)

- Unlock “resources”

unlock(lockA)

<Release> [schedule]



C++:

```
std::mutex m;  
std::lock(m);  
doCriticalwork();  
std::unlock(m);
```

# Lock

- Lock “resources”
- Process
- Unlock “resources”

<Request> [?block] <Acquired>

lock(lockA)

pred = Find(key)

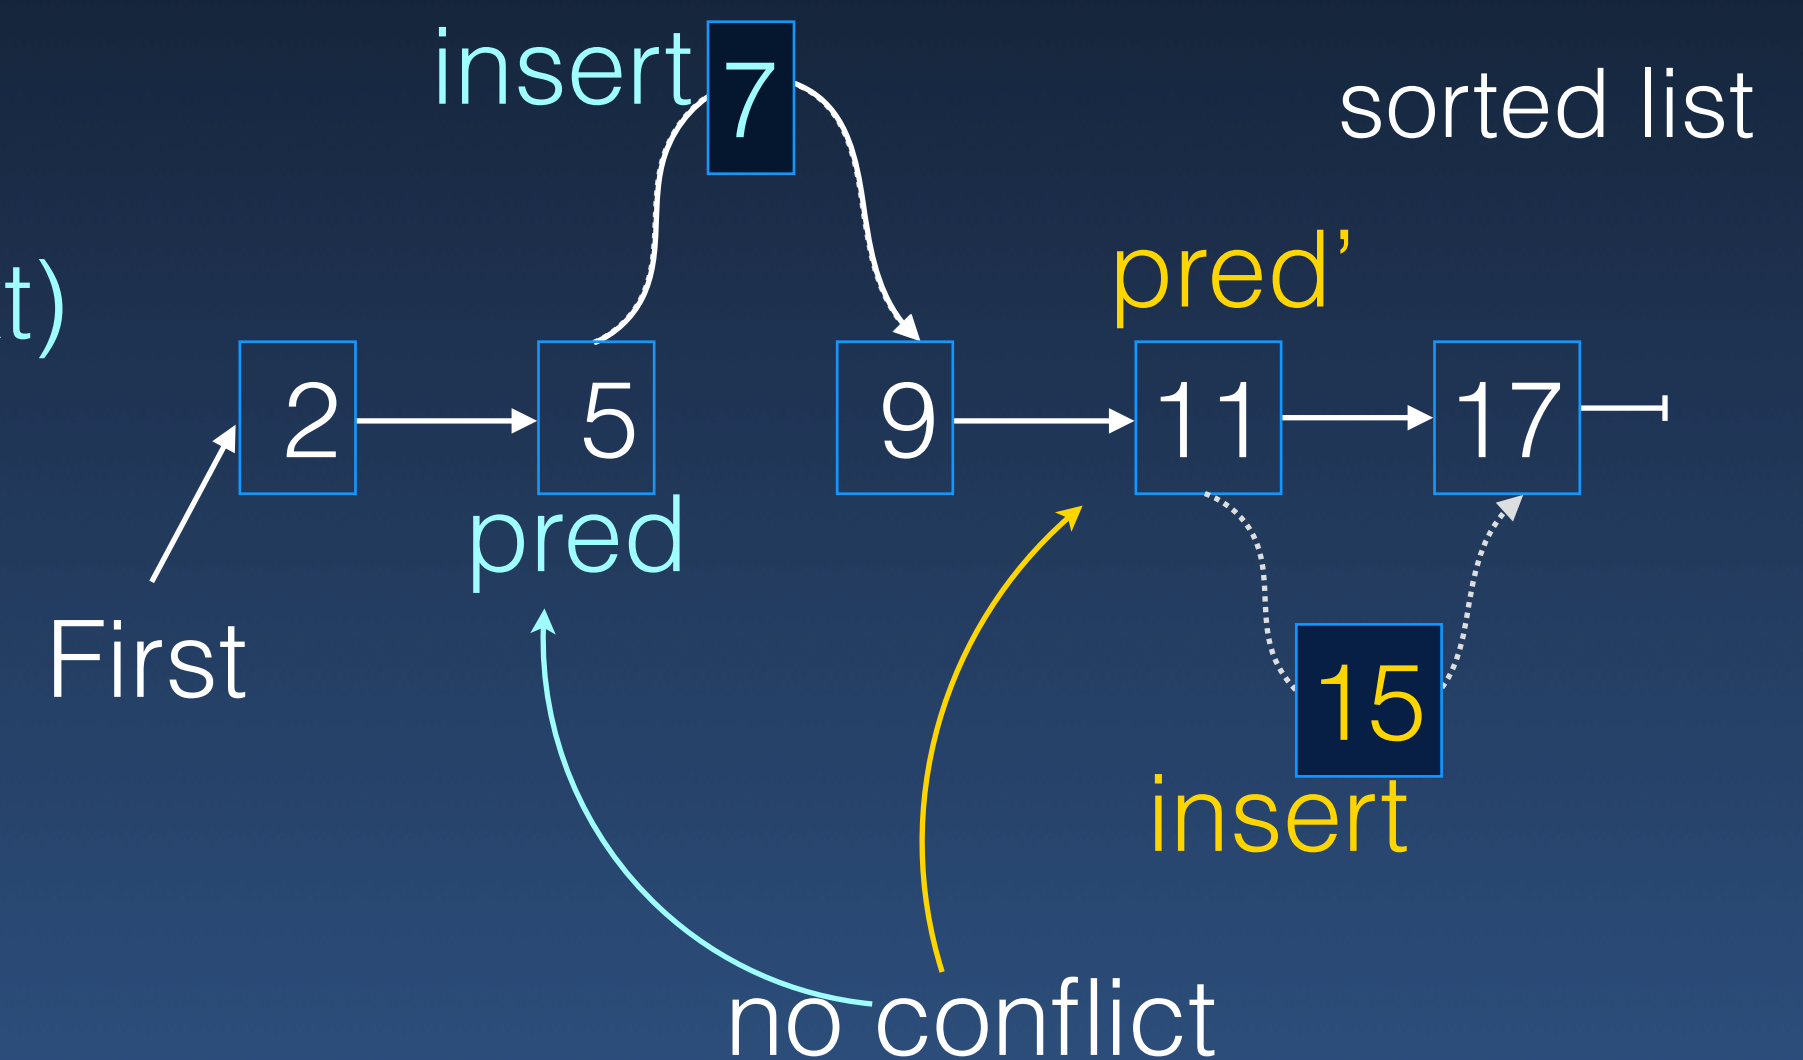
pred.nxt = Node(key, pred.nxt)

unlock(lockA)

<Release> [schedule]

C++:


```
std::mutex m;  
std::lock(m);  
doCriticalwork();  
std::unlock(m);
```



Lock the entire list?


## List using Lock

- Lock “resources”
- Process
- Unlock “resources”



## List using Lock

- Lock “resources”
- Process
- Unlock “resources”




```
Node {  
    Key   key  
    Node nxt  
    Lock  lock  
}
```



## List using Lock

- Lock “resources”
- Process
- Unlock “resources”



Insertion Loop


```
lock(pred)
if(key in [pred->key:pred->nxt->key)) {
    pred->nxt = Node(key, pred->nxt, new(Lock))
}
unlock(pred)
pred = pred->nxt
```

```
Node {
    Key   key
    Node  nxt
    Lock  lock
}
```

# List using Lock

- Lock “resources”
- Process
- Unlock “resources”

e.g.,  
`omp_set_lock(&pred->lock)`  
or, `pred->lock.lock()`




Insertion  
Loop

```
lock(pred)
if(key in [pred->key:pred->nxt->key)) {
    pred->nxt = Node(key, pred->nxt, new(Lock))
}
unlock(pred)
pred = pred->nxt
```

```
Node {
    Key   key
    Node  nxt
    Lock  lock
}
```

## List using Lock

- Lock “resources”
- Process
- Unlock “resources”



Insertion Loop


```
lock(pred)
if(key in [pred->key:pred->nxt->key)) {
    pred->nxt = Node(key, pred->nxt, new(Lock))
}
unlock(pred)
pred = pred->nxt
```

Is “9” still next to “5”?

```
Node {
    Key   key
    Node  nxt
    Lock  lock
}
```

## List using Lock

- Lock “resources”
- Process
- Unlock “resources”



Insertion  
Loop

```
lock(pred)
if(key in [pred->key:pred->nxt->key]) {
    pred->nxt = Node(key, pred->nxt, new(Lock))
}
unlock(pred)
pred = pred->nxt
```

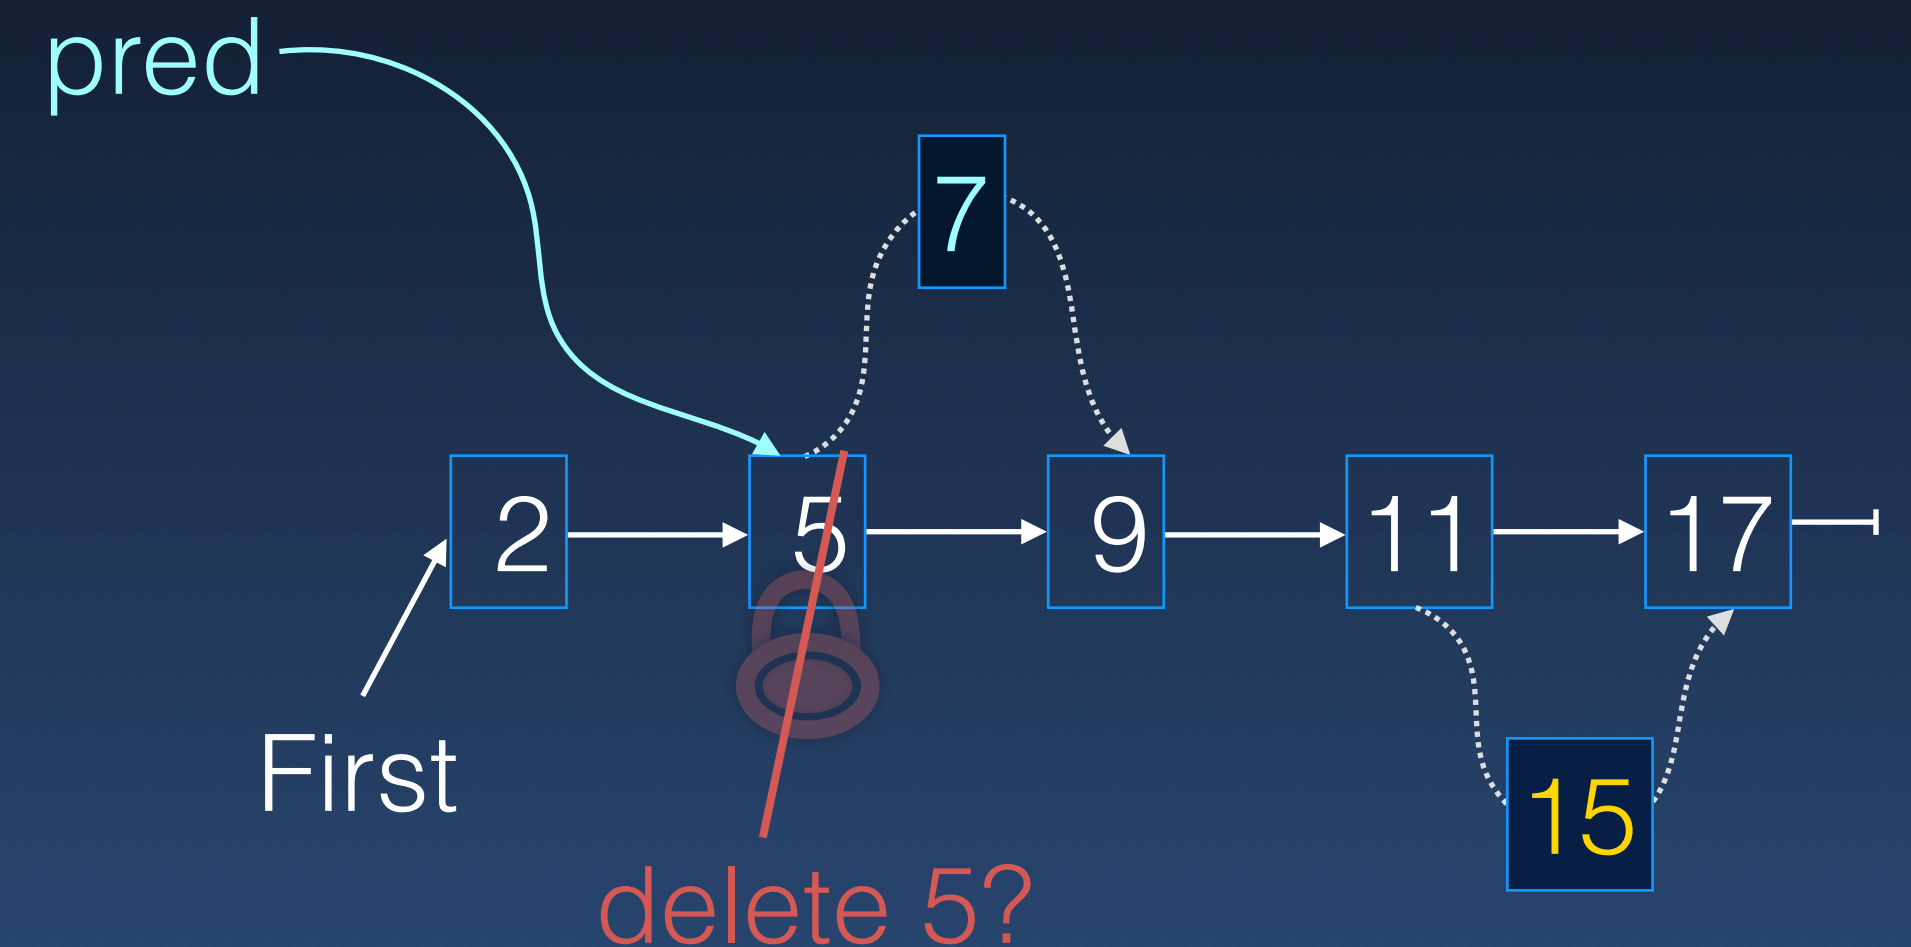
Is “9” still next to “5”?

```
Node {
    Key   key
    Node  nxt
    Lock  lock
}
```



## List using Lock

- Lock “resources”
- Process
- Unlock “resources”



Insertion Loop


```
lock(pred)
if(key in [pred->key:pred->nxt->key]) {
    pred->nxt = Node(key, pred->nxt, new(Lock))
}
unlock(pred)
pred = pred->nxt
```

Is “9” still next to “5”?

```
Node {
    Key   key
    Node  nxt
    Lock  lock
}
```

# List using Lock

- Lock “resources”
- Process
- Unlock “resources”



Insertion Loop


```
lock(pred)
if(key in [pred->key:pred->nxt->key)) {
    pred->nxt = Node(key, pred->nxt, new(Lock))
}
unlock(pred)
pred = pred->nxt
```

Before unlocking pred, capture 'nxt' locally?

```
Node {
    Key   key
    Node  nxt
    Lock  lock
}
```

## List using Lock

- Lock “resources”
- Process
- Unlock “resources”



Insertion Loop

```
lock(pred)
if(key in [pred->key:pred->nxt->key)) {
    pred->nxt = Node(key, pred->nxt, new(Lock))
}
unlock(pred)
pred = pred->nxt
```


Before unlocking pred, capture 'nxt' locally?

```
Node {
    Key   key
    Node  nxt
    Lock  lock
}
```



## List using Lock

- Lock “resources”
- Process
- Unlock “resources”



Insertion Loop

```
lock(pred)
if(key in [pred->key:pred->nxt->key)) {
    pred->nxt = Node(key, pred->nxt, new(Lock))
}
unlock(pred)
pred = pred->nxt
```

Local view?


Before unlocking pred, capture 'nxt' locally?

```
Node {
    Key   key
    Node  nxt
    Lock  lock
}
```



## List using Lock

- Lock “resources”
- Process
- Unlock “resources”



Insertion Loop


```
lock(pred) And lock(pred->nxt)
if(key in [pred->key:pred->nxt->key)) {
    pred->nxt = Node(key, pred->nxt, new(Lock))
}
unlock(pred)
pred = pred->nxt
```

Local view?

```
Node {
    Key key
    Node nxt
    Lock lock
}
```

## Insert


```
lock(First); lock(First->nxt)
pred = First; cur = pred->nxt
while(cur != Last && cur->key < key) {
    unlock(pred)
    pred = cur
    cur = cur->nxt
    lock(cur)
}
pred->nxt = new Node(key, cur, new lock())
unlock(pred); unlock(curr)
```



```
Node {
    Key   key
    Node  nxt
    Lock  lock
}
```

# Insert


```
lock(First); lock(First->nxt)
pred = First; cur = pred->nxt
while(cur != Last && cur->key < key) {
    unlock(pred)
    pred = cur
    cur = cur->nxt
    lock(cur)
}
pred->nxt = new Node(key, cur, new lock())
unlock(pred); unlock(curr)
```



```
Node {
    Key   key
    Node  nxt
    Lock  lock
}
```

# Insert

```
lock(First); lock(First->nxt)
pred = First; cur = pred->nxt
while(cur != Last && cur->key < key) {
    unlock(pred)
    pred = cur
    cur = cur->nxt
    lock(cur)
}
pred->nxt = new Node(key, cur, new lock())
unlock(pred); unlock(curr)
```




```
Node {
    Key   key
    Node  nxt
    Lock  lock
}
```



# Insert


```
lock(First); lock(First->nxt)
pred = First; cur = pred->nxt
while(cur != Last && cur->key < key) {
    unlock(pred)
    pred = cur
    cur = cur->nxt
    lock(cur)
}
pred->nxt = new Node(key, cur, new lock())
unlock(pred); unlock(curr)
```



```
Node {
    Key   key
    Node  nxt
    Lock  lock
}
```

# Delete


```
lock(First); lock(First->nxt)
pred = First; cur = pred->nxt
while(cur != Last && cur->key < key) {
    unlock(pred)
    pred = cur
    cur = cur->nxt
    lock(cur)
}
if(cur.key == key) {
    pred->nxt = cur->nxt
}
unlock(pred); unlock(curr)
```



```
Node {
    Key   key
    Node  nxt
    Lock  lock
}
```

# Lock

- Lock “resources”
- Process
- Unlock “resources”



→ Correctness depends on everyone following protocol

```
Node {  
    Key   key  
    Node nxt  
    Lock  lock  
}
```

# Peterson's Mutex Algorithm

Initially: want = {false, false}

## Thread 0

```
1 want[0] = true
2 turn = 1
3 while (want[1] && turn == 1);
4 { // critical section
  }
want[0] = false
```

## Thread 1

```
a want[1] = true
b turn = 0
c while (want[0] && turn == 0);
d { // critical section
  }
want[1] = false
```

Safe?  
(mutex guaranteed)



# Peterson's Mutex Algorithm

Initially: want = {false, false}

## Thread 0

1 want[0] = true

2 turn = 1

3

while (want[1] && turn == 1);

4

{// critical section  
}

want[0] = false

want[1]:false  
Or turn:0

## Thread 1

a want[1] = true

b turn = 0

c

while (want[0] && turn == 0);

d

{ // critical section  
}

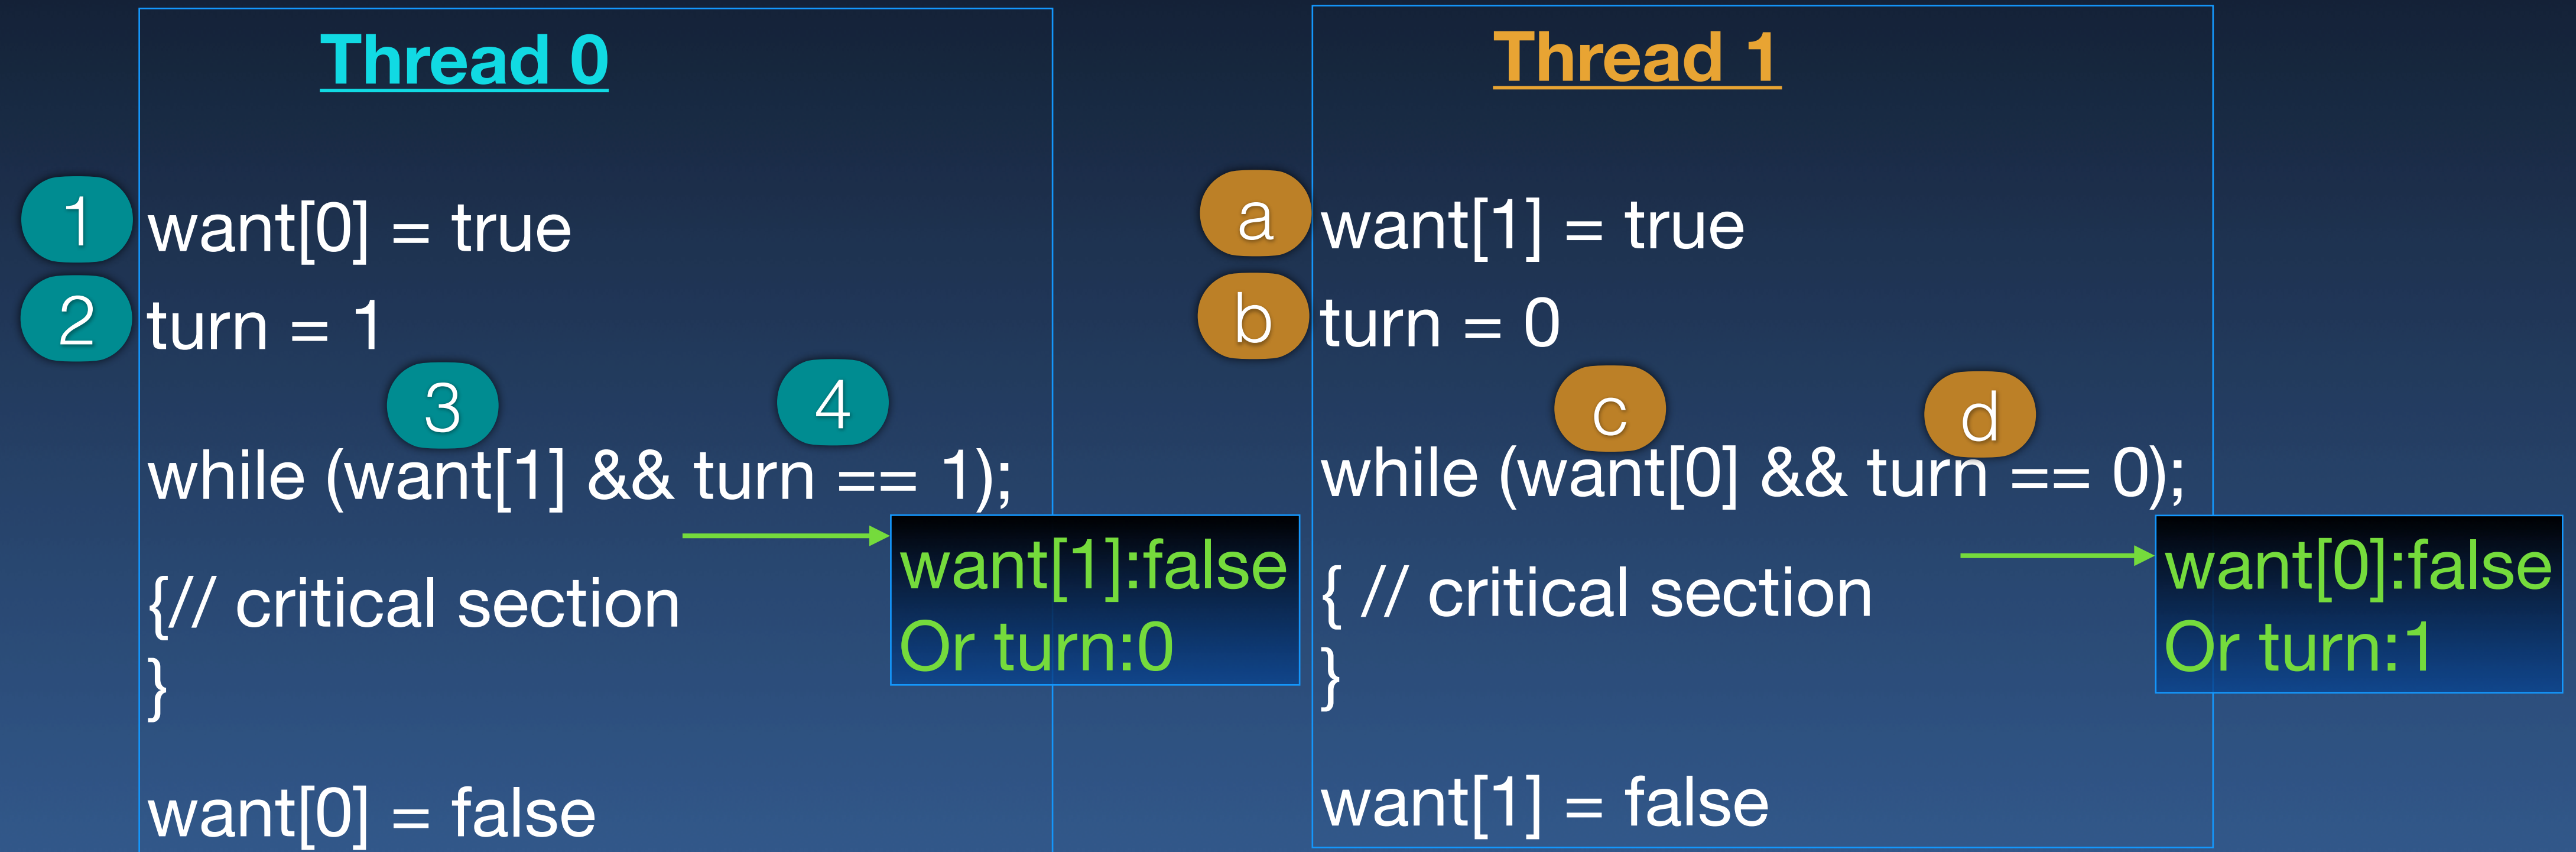
want[1] = false

want[0]:false  
Or turn:1

Safe?  
(mutex guaranteed)

# Peterson's Mutex Algorithm

Initially: want = {false, false}

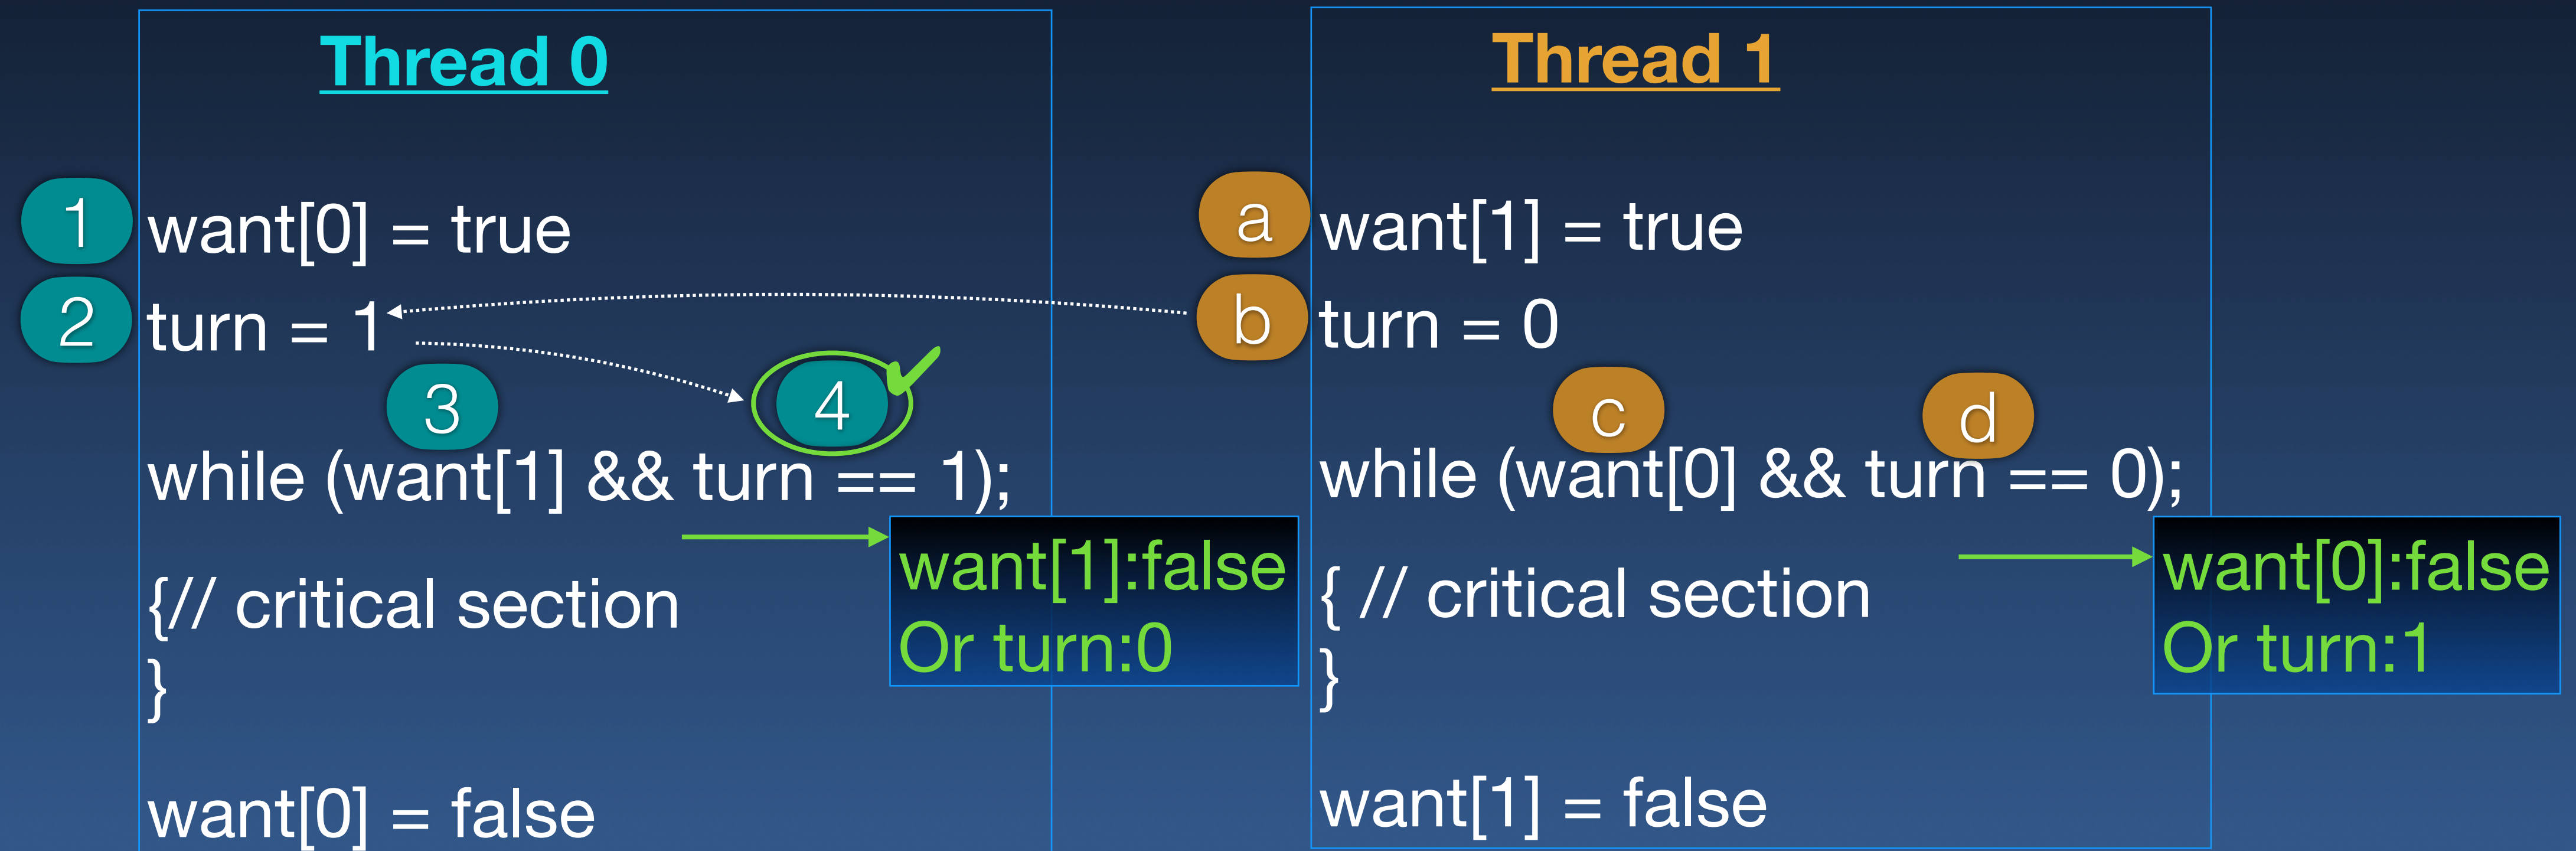


Safe?  
(mutex guaranteed)

(and Th 0 exits loop)  
Suppose: b → 2

# Peterson's Mutex Algorithm

Initially: want = {false, false}



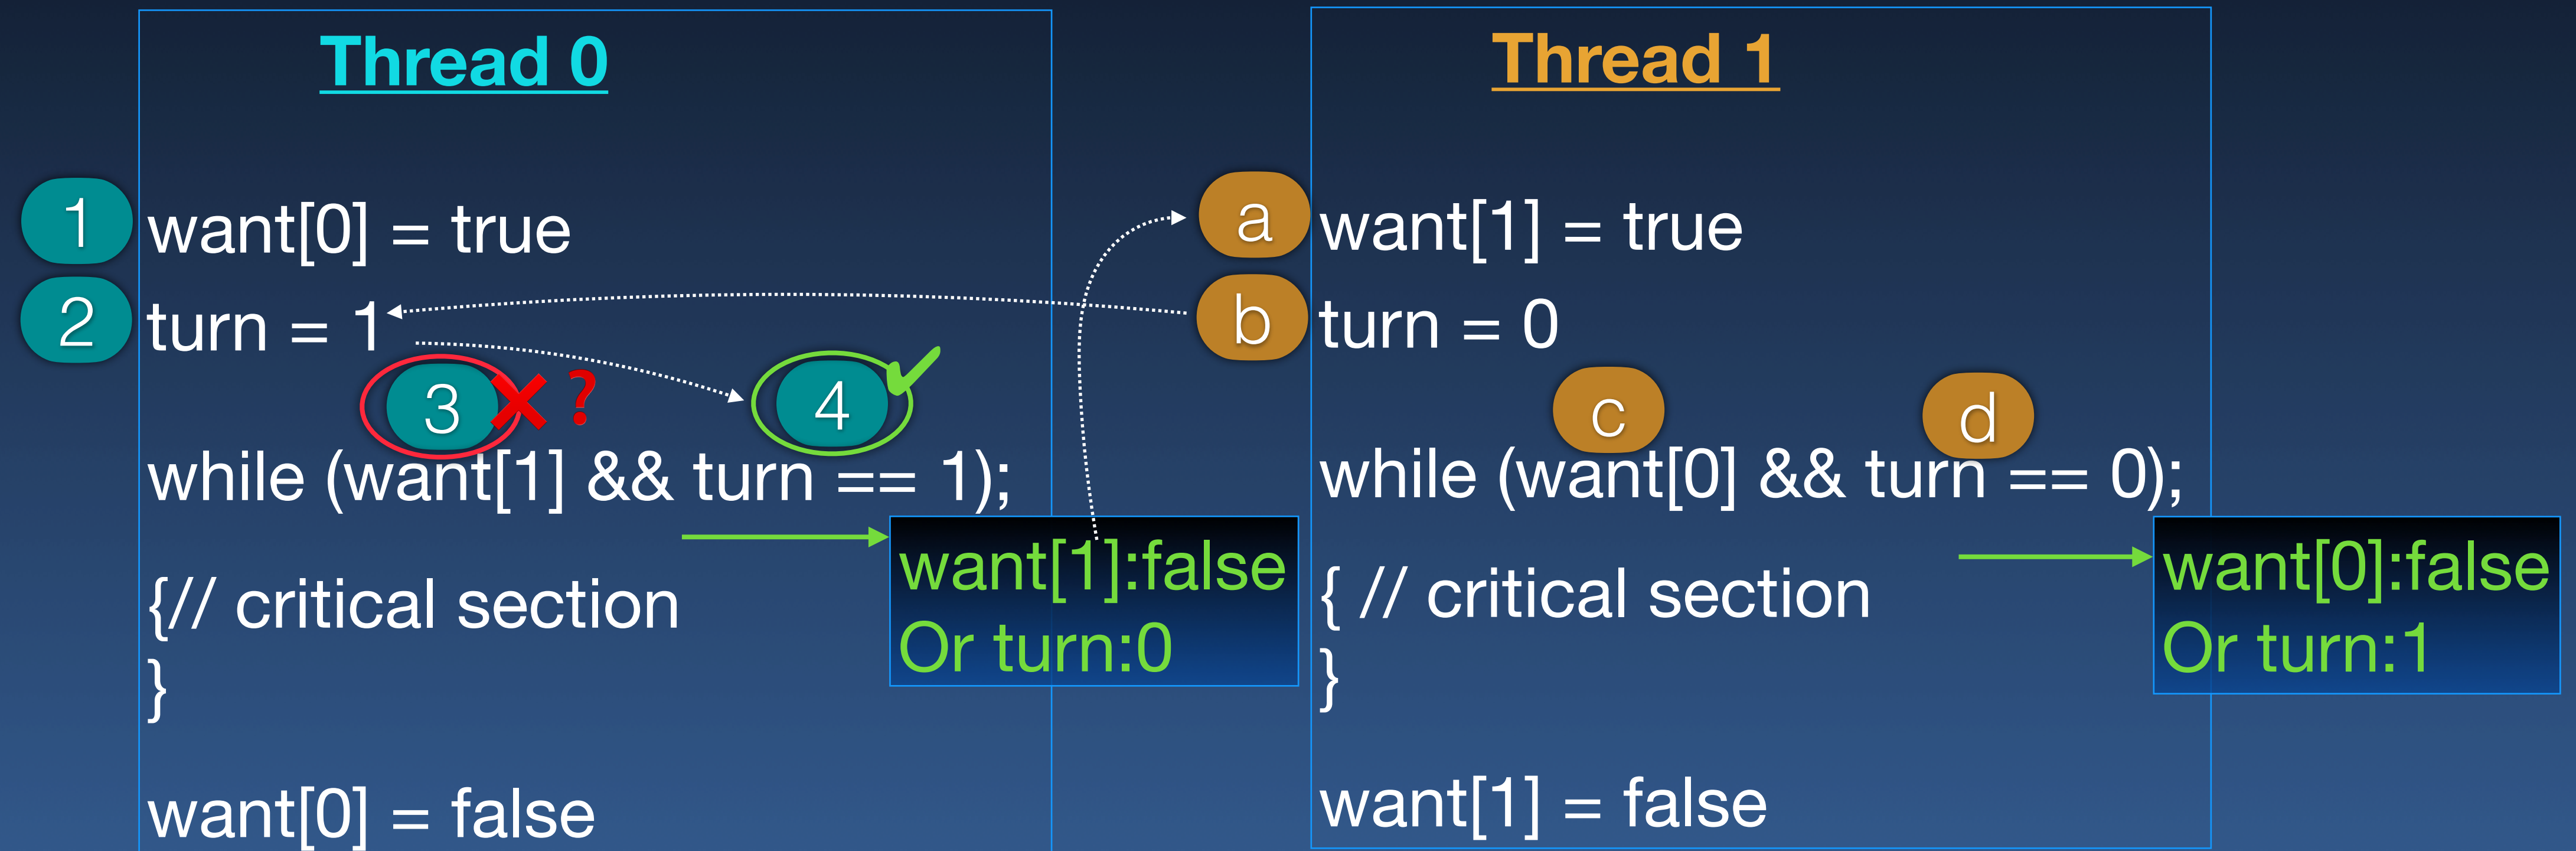
Safe?  
(mutex guaranteed)

(and Th 0 exits loop)  
Suppose: b → 2



# Peterson's Mutex Algorithm

Initially: want = {false, false}



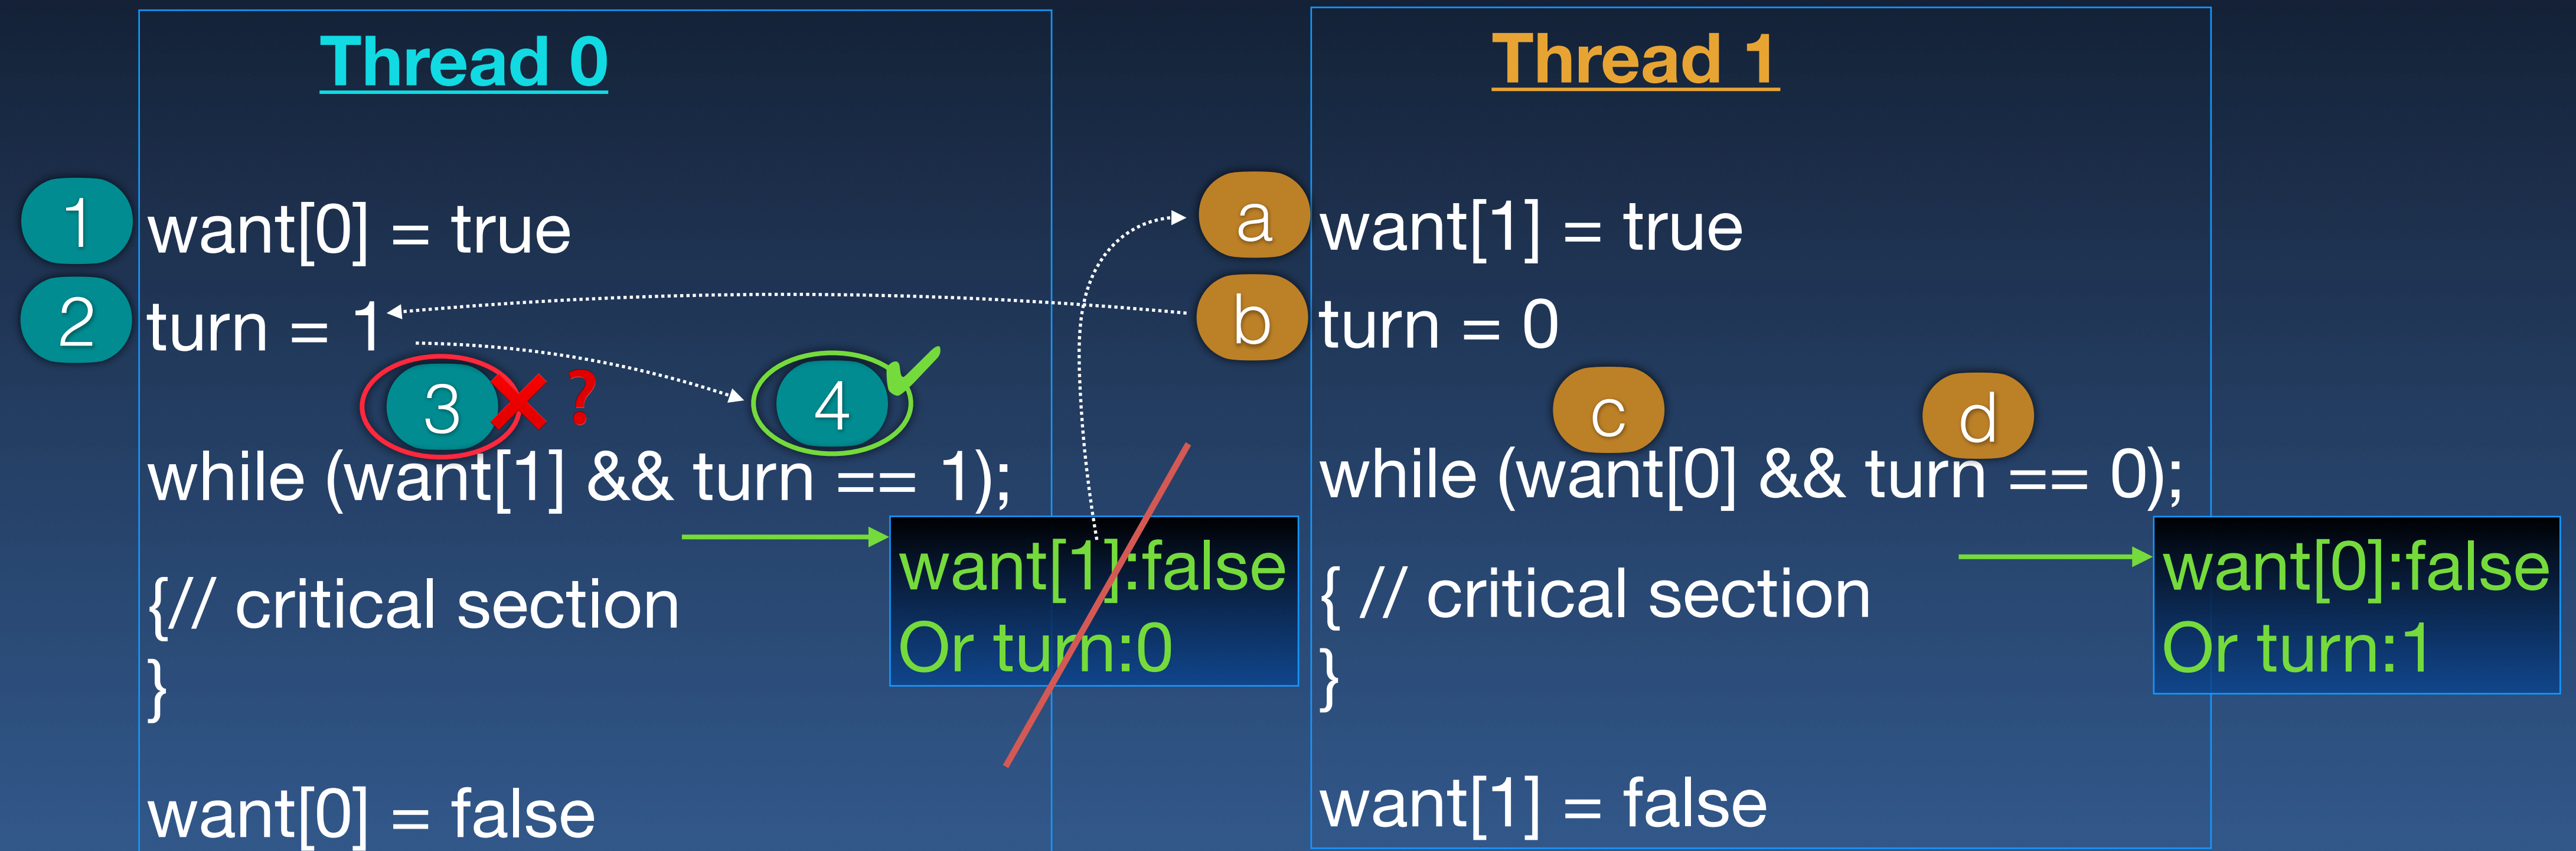
Safe?  
(mutex guaranteed)

(and Th 0 exits loop)  
Suppose: b → 2 → 3 → a



# Peterson's Mutex Algorithm

Initially: want = {false, false}



Safe?  
(mutex guaranteed)

(and Th 0 exits loop)  
Suppose: b → 2 → 3 → a

## Mutex w/Registers

- Mutual exclusion does not require hardware synchronization
- Peterson and Bakery use minimal number of registers

— Not Critical Section —

```
1: want[ID] = 1;  
2: token[ID] = 1 + max(token) ← concurrent  
3: want[ID] = 0;  
4: for other != ID {  
5:     while(want[other] == 1); ← concurrent  
6:     while(token[other] > 0 && (token[other]#other) < (token[ID]#ID);  
7: }
```

— Critical Section —

```
8: token[ID] = 0
```

Lamport's  
Bakery Algorithm

## Mutex w/Registers

- Mutual exclusion does not require hardware synchronization
- Peterson and Bakery use minimal number of registers

— Not Critical Section —

1: want[ID] = 1;

2: token[ID] = 1 + max(token) ← concurrent

3: want[ID] = 0;

4: for other != ID {

5:     while(want[other] == 1);

6:     while(token[other] > 0 && (token[other]#other) < (token[ID]#ID);

7: }

— Critical Section —

8: token[ID] = 0

Lamport's  
Bakery Algorithm

@#6, if Th.i saw token[i] < token[j], can Th.j see token[j] < token[i] (with i still in critical section)?  
If token[j] changes, it may only increase



- Mutual exclusion does not require hardware synchronization
- Peterson and Bakery use minimal number of registers
  - Too many variables? (and ever increasing counter values)

— Not Critical Section —

1: want[ID] = 1;

2: token[ID] = 1 + max(token) ← concurrent

3: want[ID] = 0;

4: for other != ID { wait if anyone in the midst of taking token

5: while(want[other] == 1); ← concurrent

6: while(token[other] > 0 && (token[other]#other) < (token[ID]#ID);

7: }


— Critical Section —

8: token[ID] = 0

Lamport's  
Bakery Algorithm

Th.i waits @#5 for Th.j to complete writing its token @#2, so that in case  $j < i$ , and both  $i$  and  $j$  get the same value, Th.i does not proceed to #6, where it may incorrectly win if write to token[j] is delayed.





✓ Lock-free? Wait-free?

```
std::atomic<node<T>*> top;  
...  
void push(const T& data) {  
    node<T>* new_node = new node<T>(data);  
    // put the current value of top into new_node->next  
    do new_node->next = top.load();  
    // make new_node the top, as long as top still equals new_node->next  
    while(!top.compare_exchange_strong(new_node->next, new_node));  
}
```

## Non-shared Logical Clock

- Each entity maintains a counter
  - ➔ increments every *step*, at its own pace
- Interaction between entities is through messages
  - ➔ Data + counter
- On message receipt:
  - ➔ If recipient counter < received count
    - ▶ Increase local counter to received count
    - ▶ Receive is also a '*step*,' so increment by one

[Lamport's Timestamp algorithm]

type      time

## Request Critical Section:

Broadcast **R** = <request, local-time>

Add **R** to local-queue

---

## Enter Critical section (**R**)

if **R.time** has the lowest time value in local-queue .AND.

Have received some **m** message from every other thread with **m.time** > **R.time**

---

## Exit Critical section (**R**):

Remove **R** from local-queue

Broadcast <release> message to all



# Distributed Mutex

## Request Critical Section:

Broadcast **R** = <request, local-time>

Add **R** to local-queue

---

## Enter Critical section (**R**)

if **R.time** has the lowest time value in local-queue .AND.

Have received some **m** message from every other thread with **m.time** > **R.time**

---

## Exit Critical section (**R**):

Remove **R** from local-queue

Broadcast <release> message to all

## Receive **R**

update(local-time)

if(**R.type** == request)

    Add **R** to local-queue

    Reply <ack, local-time>

if(**R.type** == release)

    Remove **R** from local-queue



# Distributed Mutex

## Request Critical Section:

Broadcast **R** = <request, local-time>

Add **R** to local-queue

type      time



if a request was made by another thread with time < **m.time**, it must have been received before **m**

## Enter Critical section (**R**)

if **R.time** has the lowest time value in local-queue .AND.

Have received some **m** message from every other thread with **m.time** > **R.time**

## Exit Critical section (**R**):

Remove **R** from local-queue

Broadcast <release> message to all

## Receive **R**

update(local-time)

if(**R.type** == request)

    Add **R** to local-queue

    Reply <ack, local-time>

if(**R.type** == release)

    Remove **R** from local-queue

- Synchronization primitives

- ➔ Memory fences and consistency
- ➔ Lock/Mutex, Condition variables, Atomics, Test & Set, Fetch & Add, Compare & Swap, Critical section, Barrier, Wait, Ordered
- ➔ Registers, Logical clocks, Peterson's algorithm, Bakery algorithm, Distributed mutex

- Properties of Synchronization

- ➔ Safety, Progress, Liveness
- ➔ Blocking/Non-blocking, Starvation-free, Deadlock-free, Lockfree, Waitfree